

# DBMS

## UNIT - 3

SQL

feedback/corrections: vibha@pesu.pes.edu

VIBHA MASTI  
© Vibha's notes 2021

## SQL

- Declarative language
- SQL: core components + extensions (e.g.: GIS)

### — DDL

#### 1. CREATE

- Schemas, tables, domains, views, assertions, triggers
- PSQL: database

##### 1.1 Create Schema

```
CREATE SCHEMA COMPANY AUTHORIZATION 'Jsmith';
```

- Catalog: named collection of schemas
  - Schema called INFORMATION\_SCHEMA
- Default owner: root (Postgres in PSQL)  
Default schema (creating schema): owner name

##### 1.2 Create Table

- Default schema (creating table): PUBLIC

```
CREATE TABLE EMPLOYEE
```



PUBLIC schema

- Prefix table name with Schema name while creating

```
CREATE TABLE COMPANY.EMPLOYEE
```

 schema

- Alternate: create schema name auth name (

. . .

. . .

create table statements

)

## 1.3 Virtual Relations

**CREATE VIEW NAME AS (QUERY);**

Column-name      attribute-type      constraint,  
<sub>T1</sub>

<b>CREATE TABLE EMPLOYEE</b>	( Fname Minit Lname Ssn Bdate Address Sex Salary Super_ssn Dno	VARCHAR(15) CHAR, VARCHAR(15) CHAR(9) DATE, VARCHAR(30), CHAR, DECIMAL(10,2), CHAR(9), INT	NOT NULL,  NOT NULL, NOT NULL,  NOT NULL,
	<b>PRIMARY KEY</b> (Ssn),		
<b>CREATE TABLE DEPARTMENT</b>	( Dname Dnumber Mgr_ssn Mgr_start_date	VARCHAR(15) INT CHAR(9) DATE,	NOT NULL, NOT NULL, NOT NULL,
	<b>PRIMARY KEY</b> (Dnumber),		
	<b>UNIQUE</b> (Dname),		
	<b>FOREIGN KEY</b> (Mgr_ssn) REFERENCES EMPLOYEE(Ssn) ;		
<b>CREATE TABLE DEPT_LOCATIONS</b>	( Dnumber Dlocation	INT VARCHAR(15)	NOT NULL, NOT NULL,
	<b>PRIMARY KEY</b> (Dnumber, Dlocation),		
	<b>FOREIGN KEY</b> (Dnumber) REFERENCES DEPARTMENT(Dnumber) ;		
<b>CREATE TABLE PROJECT</b>	( Pname Pnumber Plocation Dnum	VARCHAR(15) INT VARCHAR(15), INT	NOT NULL, NOT NULL,  NOT NULL,
	<b>PRIMARY KEY</b> (Pnumber),		
	<b>UNIQUE</b> (Pname),		
	<b>FOREIGN KEY</b> (Dnum) REFERENCES DEPARTMENT(Dnumber) ;		
<b>CREATE TABLE WORKS_ON</b>	( Essn Pno Hours	CHAR(9) INT DECIMAL(3,1)	NOT NULL, NOT NULL, NOT NULL,
	<b>PRIMARY KEY</b> (Essn, Pno),		
	<b>FOREIGN KEY</b> (Essn) REFERENCES EMPLOYEE(Essn),		
	<b>FOREIGN KEY</b> (Pno) REFERENCES PROJECT(Pnumber) ;		
<b>CREATE TABLE DEPENDENT</b>	( Essn Dependent_name Sex Bdate Relationship	CHAR(9) VARCHAR(15) CHAR, DATE, VARCHAR(8),	NOT NULL, NOT NULL,
	<b>PRIMARY KEY</b> (Essn, Dependent_name),		
	<b>FOREIGN KEY</b> (Essn) REFERENCES EMPLOYEE(Essn) ;		

**Figure 6.1**  
SQL CREATE  
TABLE data  
definition statements  
for defining the  
COMPANY schema  
from Figure 5.7.

## Attribute Types

### (a) String

- Fixed length: CHAR(n), CHARACTER(n)
- Variable length: VARCHAR(n), CHAR VARYING(n), CHARACTER VARYING(n)
- Enclose in '' single quotes
- Default n = 1
- Fixed length : blank-padded

### (b) Numeric

- Integer: INTEGER, INT, SMALLINT
- Float: FLOAT, REAL, DOUBLE PRECISION
- Decimal: DECIMAL(i,j), DEC(i,j), NUMERIC(i,j)
  - i: precision (total digits)
  - j: scale (no. of decimal places - default = 0)
  - DECIMAL(5,2) , DEC(10)

### (c) Bit String

- Fixed length: BIT(n), Variable length: BIT VARYING(n)
- Default n = 1
- Enclosed in '' and prefixed with B (eg: B '1001')

#### (d) Boolean

- True and False (and unknown for 3-valued logic)

#### (e) Date

- Ten positions: YYYY-MM-DD
- Enclosed in single quotes ''
- Sometimes prefixed with DATE
- Eg: DATE '2019-02-28'

#### (f) Time

- Eight positions: HH:MM:SS
- Time(i) - i extra digits for decimal of seconds

#### (g) Timestamp

- DATE + TIME + 6 positions (fraction of second)
- Optional timezone qualifier
- TIMESTAMP '2021-09-28 19:40:30.643792'

#### (h) Interval

- Relative value to increment / decrement date, time, timestamp types

additional

## Large Object Types

### (a) CLOB

- Character large object (text documents)
- PSQL: `text` (1 GB max)
- Oracle: max 4 GB
- Eg: `book review clob(10KB)`

### (b) BLOB

- Binary Large Objects
- Images, videos etc
- PSQL: `bytea` (max 1GB)
- Oracle: max 4 GB
- Eg. `image blob(10MB)`  
`movie blob(2GB)`
- PSQL : use `psycopg2` to insert images

## 1.4 Create Domain

- Renamed default datatype with or without constraints
- Easy to change datatype for domain
- Improves readability
- Eg: `CREATE DOMAIN SSN_TYPE AS CHAR(9);`

## 1.5 Create Type

- User Defined Types (UDTs)
- Object-oriented apps (not RDBMS)
- PostgreSQL: relational object DBs

`CREATE TYPE full-address AS (city VARCHAR(90), street VARCHAR(90));`

# Constraints in SQL

## 1. Attribute Constraints

### 1.1 NOT NULL

- Implicit for PK

### 1.2 DEFAULT

- DEFAULT <value>

```
vibhamasti=# CREATE TABLE STUDENT (
vibhamasti(# SID INTEGER PRIMARY KEY,
vibhamasti(# SNAME VARCHAR(2) CHECK (SNAME NOT LIKE '% %'),
vibhamasti(# CRID INTEGER,
vibhamasti(# AGE INTEGER CHECK (AGE > 20 AND AGE < 35),
vibhamasti(# DID INTEGER,
vibhamasti(# CNO INTEGER,
vibhamasti(# COLLEGE_NAME VARCHAR(20) DEFAULT 'PESU',
vibhamasti(# CONSTRAINT FK_CRID FOREIGN KEY(CRID) REFERENCES STUDENT(SID)
vibhamasti(# );
CREATE TABLE
```

```
vibhamasti=# CREATE TABLE DEPT (
vibhamasti(# DID INTEGER PRIMARY KEY,
vibhamasti(# DNAME VARCHAR(20) NOT NULL UNIQUE
vibhamasti(# );
CREATE TABLE
```

```
vibhamasti=# CREATE TABLE COURSE (
vibhamasti(# CID INTEGER PRIMARY KEY,
vibhamasti(# CNAME VARCHAR(20)
vibhamasti(# );
CREATE TABLE
```

```
vibhamasti=# ALTER TABLE STUDENT  
|ADD CONSTRAINT FK_DID FOREIGN KEY(DID) REFERENCES DEPT(DID) ON DELETE CASCADE;  
ALTER TABLE
```

```
|vibhamasti=# ALTER TABLE STUDENT DROP CONSTRAINT FK_CRID;  
ALTER TABLE  
|vibhamasti=# ALTER TABLE STUDENT ADD CONSTRAINT FK_CRID FOREIGN KEY(CRID) REFERENC  
ES STUDENT(SID) ON DELETE SET NULL;  
ALTER TABLE
```

```
vibhamasti=# ALTER TABLE STUDENT  
ADD CONSTRAINT FK_CID FOREIGN KEY(CNO) REFERENCES COURSE(CID);  
ALTER TABLE
```

```
|vibhamasti=# ALTER TABLE STUDENT ALTER COLUMN CNO SET DEFAULT 101;  
ALTER TABLE
```

```
|vibhamasti=# ALTER TABLE STUDENT ADD CONSTRAINT FK_CID FOREIGN KEY(CNO) REFERENC  
ES COURSE(CID) ON DELETE SET DEFAULT;  
ALTER TABLE
```

## 2. SELECT - FROM - WHERE

```
SELECT      <attribute list>
FROM        <table list>
WHERE       <condition>;
```

where

- <attribute list> is a list of attribute names whose values are to be retrieved by the query.
- <table list> is a list of the relation names required to process the query.
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.

### 2.2 Aliasing

```
SELECT E.Fname, S.Lname from EMPLOYEE AS E, EMPLOYEE AS S
WHERE E.SSN = S.SSN
```

### 2.3 Missing WHERE

- Select all rows
- $\pi$  relational operation
- Cartesian product - multiple tables

```
SELECT * FROM EMPLOYEE, DEPARTMENT
```

### 2.4 Joining

```
SELECT * FROM EMPLOYEE, DEPARTMENT JOIN DEPARTMENT ON (DNO=DNUMER);
```

- Can specify
  - LEFT OUTER JOIN
  - NATURAL JOIN
  - RIGHT OUTER JOIN

`SELECT * FROM EMPLOYEE, DEPARTMENT NATURAL JOIN DEPARTMENT`

## 2.5 Distinct Project

`SELECT DISTINCT ESSN FROM DEPENDENT;`

## 2.6 Default Project

- By default: select all

`SELECT ALL SALARY FROM EMPLOYEE;`

**Query 0.** Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

Q0: `SELECT Bdate, Address  
 FROM EMPLOYEE  
 WHERE Fname = 'John' AND Minit = 'B' AND Lname = 'Smith';`

**Result:**

Bdate	Address
1965-01-09	731Fondren, Houston, TX

## 2.7 Set Operations

- UNION
- EXCEPT (difference) postgres , MINUS - mysql
- INTERSECTION

## 2.8 Pattern Matching

- Postgres - only varchar
- Oracle - dates also
- LIKE
  - \_ : one char
  - % : any number of chars

SELECT FNAME, ADDRESS FROM EMPLOYEE WHERE ADDRESS LIKE '%. Houston.%';

- For dates in postgres - typecast

SELECT BDATE FROM EMPLOYEE WHERE BDATE::TEXT LIKE '1965-%';

## 2.9 BETWEEN

- Including

SELECT . . . WHERE SALARY BETWEEN 10000 AND 50000;

## 2.10 Finalized Project / Arithmetic Operations

SELECT FNAME, SALARY, 1.1\*SALARY AS HIRED FROM EMPLOYEE;

## 2.11 Ordering of Tuples

- ORDER BY
- Default: ASC
- DESC

SELECT LNAME FROM EMPLOYEE ORDER BY SALARY DESC;

- Order by multiple columns

SELECT \* FROM EMPLOYEE, DEPARTMENT JOIN DEPARTMENT ON (DNO=ONUMER)  
ORDER BY DNAME, FNAME;

- Diff orders

SELECT \* FROM EMPLOYEE, DEPARTMENT JOIN DEPARTMENT ON (DNO=ONUMER)  
ORDER BY DNAME DESC, FNAME;

SELECT \* FROM EMPLOYEE, DEPARTMENT JOIN DEPARTMENT ON (DNO=ONUMER)  
ORDER BY DNAME DESC, FNAME;

### 3. INSERT

U1:    INSERT INTO EMPLOYEE  
         VALUES ('Richard', 'K', 'Marini', '653298653', '1962-12-30', '98  
              Oak Forest, Katy, TX', 'M', 37000, '653298653', 4);

#### 3.1 Insert with Select

U3A:   CREATE TABLE WORKS\_ON\_INFO  
         ( Emp\_name           VARCHAR(15),  
          Proj\_name          VARCHAR(15),  
          Hours\_per\_week    DECIMAL(3,1) );

U3B:    INSERT INTO WORKS\_ON\_INFO ( Emp\_name, Proj\_name,  
          Hours\_per\_week )  
          SELECT E.Lname, P.Pname, W.Hours  
          FROM PROJECT P, WORKS\_ON W, EMPLOYEE E  
          WHERE P.Pnumber = W.Pno AND W.Essn = E.Ssn;

## 4. DELETE

U4A:	DELETE FROM	EMPLOYEE
	WHERE	Lname = 'Brown';
U4B:	DELETE FROM	EMPLOYEE
	WHERE	Ssn = '123456789';
U4C:	DELETE FROM	EMPLOYEE
	WHERE	Dno = 5;
U4D:	DELETE FROM	EMPLOYEE;

## 5. UPDATE

U5:	UPDATE	PROJECT
	SET	Plocation = 'Bellaire', Dnum = 5
	WHERE	Pnumber = 10;

## ADVANCED DATATYPES — LARGE OBJECTS

- Postgres:
  1. text: large text (Oracle - clob)
  2. bytea: images (byte array) (Oracle - blob)
  3. ntext: even larger text
- psycopg2: python

## COMPLEX SQL

### 1.1 NULL & Three-Valued Logic

- NULL due to 3 reasons

**1. Unknown value.** A person's date of birth is not known, so it is represented by NULL in the database. An example of the other case of unknown would be NULL for a person's home phone because it is not known whether or not the person has a home phone.

**2. Unavailable or withheld value.** A person has a home phone but does not want it to be listed, so it is withheld and represented as NULL in the database.

**3. Not applicable attribute.** An attribute LastCollegeDegree would be NULL for a person who has no college degrees because it does not apply to that person.

- Three-valued logic (not bool) — True, False, Unknown

**Table 7.1** Logical Connectives in Three-Valued Logic

		TRUE	FALSE	UNKNOWN
(a) AND		TRUE	FALSE	UNKNOWN
		TRUE	FALSE	UNKNOWN
(b) OR	TRUE	TRUE	TRUE	TRUE
	FALSE	TRUE	FALSE	UNKNOWN
	UNKNOWN	TRUE	UNKNOWN	UNKNOWN
(c) NOT	TRUE	FALSE		
	FALSE	TRUE		
	UNKNOWN	UNKNOWN		

Eg

**Query 18.** Retrieve the names of all employees who do not have supervisors.

**Q18:**

```
SELECT      Fname, Lname
FROM        EMPLOYEE
WHERE       Super_ssn IS NULL;
```

## 1.2 Nested Queries, Tuples and set/ Multiset Comparisons

**Query 4.** Make a list of all project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.

**Q4A:**

```
( SELECT    DISTINCT Pnumber
  FROM      PROJECT, DEPARTMENT, EMPLOYEE
  WHERE     Dnum = Dnumber AND Mgr_ssn = Ssn
            AND     Lname = 'Smith'
UNION
( SELECT    DISTINCT Pnumber
  FROM      PROJECT, WORKS_ON, EMPLOYEE
  WHERE     Pnumber = Pno AND Essn = Ssn
            AND     Lname = 'Smith' );
```

using nested

**Q4A:**

```
SELECT    DISTINCT Pnumber
  FROM      PROJECT
  WHERE     Pnumber IN
( SELECT    Pnumber
  FROM      PROJECT, DEPARTMENT, EMPLOYEE
  WHERE     Dnum = Dnumber AND
            Mgr_ssn = Ssn AND Lname = 'Smith' )
OR
Pnumber IN
( SELECT    Pno
  FROM      WORKS_ON, EMPLOYEE
  WHERE     Essn = Ssn AND Lname = 'Smith' );
```

## Tuple Comparisons

### 1.2.1 IN

... where (essn, pno) in (select essn, pno ... where);

```
SELECT      DISTINCT Essn
FROM        WORKS_ON
WHERE       (Pno, Hours) IN
            (SELECT      Pno, Hours
             FROM        WORKS_ON
             WHERE       Essn = '123456789');
```

• IN operator is same as = ANY

### 1.2.2 ANY & ALL

• salary > all sals in dept 5

```
SELECT      Lname, Fname
FROM        EMPLOYEE
WHERE       Salary > ALL
            (SELECT      Salary
             FROM        EMPLOYEE
             WHERE       Dno = 5);
```

• = all & > all may give  $\emptyset$  in most instances  
• = any or = some

### 1.2.3 ALIAS

**Query 16.** Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.

**Q16:**

```
SELECT      E.Fname, E.Lname
FROM        EMPLOYEE AS E
WHERE       E.Ssn IN
            (SELECT      D.Essn
             FROM        DEPENDENT AS D
             WHERE       E.Fname = D.Dependent_name
                         AND E.Sex = D.Sex);
```

- PSQL: no need to write AS (optional)

## 1.3 CORRELATED QUERIES

- correlated queries: for every tuple in outer query, inner query is executed (inner query uses result of outer query)
- normal: inner then outer

Q16A:    SELECT                E.Fname, E.Lname  
           FROM                 EMPLOYEE AS E, DEPENDENT AS D  
           WHERE                E.Ssn = D.Essn **AND** E.Sex = D.Sex  
                               AND E.Fname = D.Dependent\_name;

✓ more optimised

## 1.4 EXISTS and UNIQUE

- check if inner query is empty

Q16B:    SELECT                E.Fname, E.Lname  
           FROM                 EMPLOYEE AS E  
           WHERE                **EXISTS** ( SELECT                \*  
                                FROM                 DEPENDENT AS D  
                                WHERE                E.Ssn = D.Essn **AND** E.Sex = D.Sex  
                               AND E.Fname = D.Dependent\_name);

like IN



Query 6. Retrieve the names of employees who have no dependents.

Q6:    SELECT                Fname, Lname  
           FROM                 EMPLOYEE  
           WHERE                **NOT EXISTS** ( SELECT                \*  
                                FROM                 DEPENDENT  
                                WHERE                Ssn = Essn );

(a)

select fname, lname from employee except

(select fname, lname from employee, dependent where essn=ssn)

(b)

Select fname, lname from employee where ssn not in (select  
essn from dependent)

Eg

Query 7. List the names of managers who have at least one dependent.

Q7:    **SELECT**      Fname, Lname  
         **FROM**        EMPLOYEE  
         **WHERE**     **EXISTS** ( **SELECT**      \*  
                        **FROM**        DEPENDENT  
                        **WHERE**      Ssn = Essn )  
         **AND**  
         **EXISTS** ( **SELECT**      \*  
                        **FROM**        DEPARTMENT  
                        **WHERE**      Ssn = Mgr\_ssn );

(a)

select distinct fname, lname from employee , department,  
dependent where ssn=mgr\_ssn and ssn=essn;

### 1.4.1 For all / division is SQL

- Using EXISTS and NOT EXISTS

Q3: Retrieve the name of each employee who works on all the projects controlled by department number 5

Q3A: `SELECT Fname, Lname  
FROM EMPLOYEE  
WHERE NOT EXISTS (( SELECT Pnumber  
FROM PROJECT  
WHERE Dnum = 5)  
EXCEPT ( SELECT Pno  
WORKS_ON  
WHERE Ssn = Essn ));`

*leftover projects in dept 5 for emp*

*should not have leftover*

*all projects emp works in*

Q3B: `SELECT Lname, Fname  
FROM EMPLOYEE  
WHERE NOT EXISTS ( SELECT *  
FROM WORKS_ON B  
WHERE ( B.Pno IN ( SELECT Pnumber  
FROM PROJECT  
WHERE Dnum = 5 )  
AND  
NOT EXISTS ( SELECT *  
FROM WORKS_ON C  
WHERE C.Essn = Ssn  
AND C.Pno = B.Pno ))));`

## 1.5 Explicit Set of values in where

Query 17. Retrieve the Social Security numbers of all employees who work on project numbers 1, 2, or 3.

Q17:    **SELECT**    **DISTINCT** Essn  
        **FROM**        WORKS\_ON  
        **WHERE**      Pno IN (1, 2, 3);

## 1.6 JOIN in FROM

Q1A:    **SELECT**    Fname, Lname, Address  
        **FROM**        (EMPLOYEE JOIN DEPARTMENT ON Dno = Dnumber)  
        **WHERE**      Dname = 'Research';

- here also NULL ignored

### 1.6.1 Natural Join

Q1B:    **SELECT**    Fname, Lname, Address  
        **FROM**        (EMPLOYEE NATURAL JOIN  
                      (DEPARTMENT AS DEPT (Dname, Dno, Mssn, Msdate)))  
        **WHERE**      Dname = 'Research';

- gives cartesian prod if col names do not match

### 1.6.2 Outer Joins

Q8B:    **SELECT**    E.Lname **AS** Employee\_name,  
          S.Lname **AS** Supervisor\_name  
        **FROM**        (EMPLOYEE **AS** E LEFT OUTER JOIN EMPLOYEE **AS** S  
                  ON E.Super\_ssn = S.Ssn);

shows emps w/o super  
also

## Alternate syntax (Oracle only)

```
Q8C:   SELECT E.Lname, S.Lname  
        FROM EMPLOYEE E, EMPLOYEE S  
       WHERE E.Super_ssn += S.Ssn;
```

### 1.7 Aggregate Functions

**Query 19.** Find the sum of the salaries of all employees, the maximum salary, the minimum salary, and the average salary.

```
Q19:   SELECT SUM (Salary), MAX (Salary), MIN (Salary), AVG (Salary)  
        FROM EMPLOYEE;
```

- NULL ignored from aggregate func on particular col
- count(\*) includes NULL in few columns
- Max, min, sum, avg, count

**Query 20.** Find the sum of the salaries of all employees of the 'Research' department, as well as the maximum salary, the minimum salary, and the average salary in this department.

```
Q20:   SELECT SUM (Salary), MAX (Salary), MIN (Salary), AVG (Salary)  
        FROM (EMPLOYEE JOIN DEPARTMENT ON Dno = Dnumber)  
       WHERE Dname = 'Research';
```

- instead of where use having if group by used (filter out groups)

**Queries 21 and 22.** Retrieve the total number of employees in the company (Q21) and the number of employees in the 'Research' department (Q22).

```
Q21:   SELECT COUNT (*)  
        FROM EMPLOYEE;
```

```
Q22:   SELECT COUNT (*)  
        FROM EMPLOYEE, DEPARTMENT  
       WHERE DNO = DNUMBER AND DNAME = 'Research';
```

- select clause in group by is very imp (project)
- valid if 1 value per group

**Query 24.** For each department, retrieve the department number, the number of employees in the department, and their average salary.

**Q24:**    **SELECT**      Dno, COUNT (\*), AVG (Salary)  
**FROM**                    EMPLOYEE  
**GROUP BY**      Dno;

(a)

Fname	Minit	Lname	Ssn	...	Salary	Super_ssn	Dno	
John	B	Smith	123456789		30000	333445555	5	
Franklin	T	Wong	333445555		40000	888665555	5	
Ramesh	K	Narayan	666884444		38000	333445555	5	
Joyce	A	English	453453453	...	25000	333445555	5	
Alicia	J	Zelaya	999887777		25000	987654321	4	
Jennifer	S	Wallace	987654321		43000	888665555	4	
Ahmad	V	Jabbar	987987987		25000	987654321	4	
James	E	Bong	888665555		55000	NULL	1	

Dno	Count (*)	Avg (Salary)
5	4	33250
4	3	31000
1	1	55000

Result of Q24

Grouping EMPLOYEE tuples by the value of Dno

- Group by — NULL is one group
- Group by with having

**Query 26.** For each project *on which more than two employees work*, retrieve the project number, the project name, and the number of employees who work on the project.

**Q26:**    **SELECT**      Pnumber, Pname, COUNT (\*)  
**FROM**                    PROJECT, WORKS\_ON  
**WHERE**                  Pnumber = Pno  
**GROUP BY**      Pnumber, Pname  
**HAVING**               COUNT (\*) > 2;

## Result

(b)

Pname	Pnumber	...	Essn	Pno	Hours
ProductX	1	...	123456789	1	32.5
ProductX	1		453453453	1	20.0
ProductY	2		123456789	2	7.5
ProductY	2		453453453	2	20.0
ProductY	2		333445555	2	10.0
ProductZ	3	...	666884444	3	40.0
ProductZ	3		333445555	3	10.0
Computerization	10		333445555	10	10.0
Computerization	10		999887777	10	10.0
Computerization	10		987987987	10	35.0
Reorganization	20		333445555	20	10.0
Reorganization	20		987654321	20	15.0
Reorganization	20		888665555	20	NULL
Newbenefits	30		987987987	30	5.0
Newbenefits	30		987654321	30	20.0
Newbenefits	30		999887777	30	30.0

These groups are not selected by the HAVING condition of Q26.

After applying the WHERE clause but before applying HAVING

Pname	Pnumber	...	Essn	Pno	Hours	Pname	Count (*)
ProductY	2	...	123456789	2	7.5	ProductY	3
ProductY	2		453453453	2	20.0	Computerization	3
ProductY	2		333445555	2	10.0	Reorganization	3
Computerization	10	...	333445555	10	10.0	Newbenefits	3
Computerization	10		999887777	10	10.0		
Computerization	10		987987987	10	35.0		
Reorganization	20		333445555	20	10.0		
Reorganization	20		987654321	20	15.0		
Reorganization	20		888665555	20	NULL		
Newbenefits	30		987987987	30	5.0		
Newbenefits	30		987654321	30	20.0		
Newbenefits	30		999887777	30	30.0		

Result of Q26  
(Pnumber not shown)

After applying the HAVING clause condition

- Group by multiple

**Query 27.** For each project, retrieve the project number, the project name, and the number of employees from department 5 who work on the project.

**Q27:**

```
SELECT      Pnumber, Pname, COUNT (*)
FROM        PROJECT, WORKS_ON, EMPLOYEE
WHERE       Pnumber = Pno AND Ssn = Essn AND Dno = 5
GROUP BY    Pnumber, Pname;
```

**Q:** Count the total number of employees whose salaries exceed \$40,000 in each department, but only for departments where more than five employees work

incorrect:

execution  
order

```
SELECT      Dno, COUNT (*)
FROM        EMPLOYEE
WHERE       Salary > 40000
GROUP BY    Dno
HAVING     COUNT (*) > 5;
```

will first filter emp  
before counting dept  
size

correct:

**Q28:**

```
SELECT      Dno, COUNT (*)
FROM        EMPLOYEE
WHERE       Salary > 40000 AND Dno IN
           ( SELECT      Dno
             FROM        EMPLOYEE
             HAVING     COUNT (*) > 5 )
GROUP BY    Dno;
```

get dnos

{

}

## 1.8 With and Case clauses

- **WITH:** common table expressions (CTE)
  - write auxiliary statements
  - define temp table for one query

eg: with abc as optional keyword  
(  
    complex query  
    :  
)  
select \* from abc;

- **WITH** mainly for readability ; can use nested query instead

Q28': WITH BIGDEPTS (Dno) AS  
( SELECT Dno  
    FROM EMPLOYEE  
    GROUP BY Dno  
    HAVING COUNT (\*) > 5)  
SELECT Dno, COUNT (\*)  
FROM EMPLOYEE  
WHERE Salary > 40000 AND Dno IN BIGDEPTS  
GROUP BY Dno;

- **CASE:** like switch case

U6': UPDATE EMPLOYEE  
SET Salary =  
CASE  
    WHEN Dno = 5 THEN Salary + 2000  
    WHEN Dno = 4 THEN Salary + 1500  
    WHEN Dno = 1 THEN Salary + 3000  
    ELSE Salary + 0 ;

## Assertions & Triggers

## 2.1 ASSERTIONS

- Assertions: constraints outside scope of built-in relational model constraints

```
CREATE ASSERTION SALARY_CONSTRAINT
CHECK ( NOT EXISTS ( SELECT      *
                      FROM EMPLOYEE E, EMPLOYEE M,
                           DEPARTMENT D
                     WHERE E.Salary > M.Salary
                           AND E.Dno = D.Dnumber
                           AND D.Mgr_ssN = M.SsN ) );
```

- Not in pg, only Oracle, MySQL

## 2.2 TRIGGERS

- Triggers: actions taken when event occurs
  - Components: event(s), condition, action (CECA rule)
  - Before or after event

```
R5: CREATE TRIGGER SALARY_VIOLATION  
    BEFORE INSERT OR UPDATE OF SALARY, SUPERVISOR_SSN  
    ON EMPLOYEE  
    FOR EACH ROW  
    WHEN ( NEW.SALARY > ( SELECT SALARY FROM EMPLOYEE  
                            WHERE SSN = NEW.SUPERVISOR_SSN ) )  
        INFORM_SUPERVISOR(NEW.Supervisor_ssn,  
                           NEW.Ssn);
```

 action (function/procedure)  
can return no ret

## Functions

- $\backslash df \rightarrow$  list
- Function: 0 or more params
- Procedural language for psql : plpgsql (we use)
  - can use others
- CREATE [OR REPLACE] FUNCTION name(params)  
RETURNS type  
LANGUAGE plpgsql  
AS

```
$$
declare
-- variable declaration (local)
begin
-- logic
end;
$$
```

```
vibhamasti=# create table accounts (
vibhamasti(# id int generated by default as identity,
vibhamasti(# name varchar(100) not null,
vibhamasti(# balance dec(15, 2) not null,
vibhamasti(# primary key(id)
vibhamasti(# );
CREATE TABLE
vibhamasti=# insert into accounts(name, balance) values ('Bob', 10000);
INSERT 0 1
vibhamasti=# insert into accounts(name, balance) values ('Alice', 10000);
INSERT 0 1
vibhamasti=# insert into accounts(name, balance) values ('Colin', 10000);
INSERT 0 1
```

```
vibhamasti=# create function get_total(from_id int, to_id int)
returns int
language plpgsql
as
$$
declare
balance_total integer;
begin
select sum(balance) into balance_total from accounts where id between from_id an
d to_id; return balance_total;
end;
$$;
CREATE FUNCTION
```

```
[vibhamasti=# select * from get_total(1, 2);
get_total
-----
20000
(1 row)
```

## Trigger

CREATE TRIGGER \_\_\_\_\_  
BEFORE/AFTER \_\_\_\_\_  
ON \_\_\_\_\_

FOR EACH ROW EXECUTE PROCEDURE

CREATE PROCEDURE
no return
cursor.call()

call procedure cur.execute("CALL procedure(%s,%s)",( ))
cur.callproc ('func\_name', (val1, val2))

## VIEW

- Virtual table

V1:	CREATE VIEW AS SELECT FROM WHERE	WORKS_ON1 Fname, Lname, Pname, Hours EMPLOYEE, PROJECT, WORKS_ON Ssn = Essn <b>AND</b> Pno = Pnumber;
V2:	CREATE VIEW AS SELECT FROM WHERE GROUP BY	DEPT_INFO(Dept_name, No_of_emps, Total_sal) Dname, <b>COUNT (*)</b> , <b>SUM</b> (Salary) DEPARTMENT, EMPLOYEE Dnumber = Dno Dname;

- Can make a view of a join

## Strategies

### (a) Query modification approach

- Maps view name to query (sub-query)

QV1:    **SELECT**      Fname, Lname  
**FROM**      WORKS\_ON1  
**WHERE**     Pname = 'ProductX';

↓  
maps

**SELECT**      Fname, Lname  
**FROM**      EMPLOYEE, PROJECT, WORKS\_ON  
**WHERE**     Ssn = Essn **AND** Pno = Pnumber  
**AND** Pname = 'ProductX';

### (b) View Materialisation / Realisation

- Stores view physically
- Updation strategies

- immediate update
- lazy update (most)
- periodic update

## Updation of Views

- A view with a single defining table is updatable if the view attributes contain the primary key of the base relation, as well as all attributes with the NOT NULL constraint that do not have default values specified.
- Views defined on multiple tables using joins are generally not updatable.
- Views defined using grouping and aggregate functions are not updatable.

UPDATE < > SET < > WHERE < >

## View Constraints

CREATE VIEW < > AS < > WITH CHECK OPTION;

- Fails row constraint otherwise

## View as an Authorisation Mechanism

- Restrict access to hidden details
- Grant permissions to users

**Table 7.2** Summary of SQL Syntax

CREATE TABLE <table name> ( <column name><column type> [ <attribute constraint> ]  
  { ,<column name><column type> [ <attribute constraint> ] }  
  [ <table constraint> { ,<table constraint> } ] )

---

DROP TABLE <table name>

---

ALTER TABLE <table name> ADD <column name> <column type>

---

SELECT [ DISTINCT ] <attribute list>  
FROM ( <table name> { <alias> } | <joined table> ) { , ( <table name> { <alias> } | <joined table> ) }  
[ WHERE <condition> ]  
[ GROUP BY <grouping attributes> [ HAVING <group selection condition> ] ]  
[ ORDER BY <column name> [ <order> ] { ,<column name> [ <order> ] } ]

---

<attribute list> ::= ( \* | ( <column name> | <function> ( ( [ DISTINCT ] <column name> | \* ) )  
  { , ( <column name> | <function> ( ( [ DISTINCT ] <column name> | \* ) ) ) ) ) )

---

<grouping attributes> ::= <column name> { , <column name> }

---

<order> ::= ( ASC | DESC )

---

INSERT INTO <table name> [ ( <column name> { , <column name> } ) ]  
( VALUES ( <constant value> , { <constant value> } ) { , ( <constant value> { , <constant value> } ) }  
| <select statement> )

---

DELETE FROM <table name>  
[ WHERE <selection condition> ]

---

UPDATE <table name>  
SET <column name> = <value expression> { , <column name> = <value expression> }  
[ WHERE <selection condition> ]

---

CREATE [ UNIQUE] INDEX <index name>  
ON <table name> ( <column name> [ <order> ] { , <column name> [ <order> ] } )  
[ CLUSTER ]

---

DROP INDEX <index name>

---

CREATE VIEW <view name> [ ( <column name> { , <column name> } ) ]  
AS <select statement>

---

DROP VIEW <view name>

NOTE: The commands for creating and dropping indexes are not part of standard SQL.

## OLAP

- Online Analytical Processing
- Allows data to be summarised and viewed in different ways in an online fashion
- Negligible delay
- For explanation: schema  
`sales (item_name, color, clothes_size, quantity)`
- **Measure attributes:** measure some value that can be aggregated upon
  - eg: attribute quantity of the sales relation
- **Dimension attributes:** define the dimensions on which measure attributes are viewed
  - eg: attributes item-name, color and size of the sales relation
- Data that can be modelled as dimension attributes and measure attributes are called **multidimensional data**

Suppose that *item\_name* can take on the values (skirt, dress, shirt, pants), *color* can take on the values (dark, pastel, white), *clothes\_size* can take on values (small, medium, large), and *quantity* is an integer value representing the total number of items of a given *{item\_name, color, clothes\_size}*. An instance of the *sales* relation is shown in Figure 5.16.

item_name	color	clothes_size	quantity
skirt	dark	small	2
skirt	dark	medium	5
skirt	dark	large	1
skirt	pastel	small	11
skirt	pastel	medium	9
skirt	pastel	large	15
skirt	white	small	2
skirt	white	medium	5
skirt	white	large	3
dress	dark	small	2
dress	dark	medium	6
dress	dark	large	12
dress	pastel	small	4
dress	pastel	medium	3
dress	pastel	large	3
dress	white	small	2
dress	white	medium	3
dress	white	large	0
shirt	dark	small	2
shirt	dark	medium	6
shirt	dark	large	6
shirt	pastel	small	4
shirt	pastel	medium	1
shirt	pastel	large	2
shirt	white	small	17
shirt	white	medium	1
shirt	white	large	10
pants	dark	small	14
pants	dark	medium	6
pants	dark	large	0
pants	pastel	small	1
pants	pastel	medium	0
pants	pastel	large	1
pants	white	small	3
pants	white	medium	0
pants	white	large	2

Figure 5.16 An example of *sales* relation.

item_name	color	clothes_size	quantity
skirt	dark	all	8
skirt	pastel	all	35
skirt	white	all	10
skirt	all	all	53
dress	dark	all	20
dress	pastel	all	10
dress	white	all	5
dress	all	all	35
shirt	dark	all	14
shirt	pastel	all	7
shirt	white	all	28
shirt	all	all	49
pants	dark	all	20
pants	pastel	all	2
pants	white	all	5
pants	all	all	27
all	dark	all	62
all	pastel	all	54
all	white	all	48
all	all	all	164

Figure 5.21 Relational representation of the data in Figure 5.17.

## L. CROSS-TABULATION / PIVOT TABLE

- Row headers — values of one attr
- Column headers — values of another attr

		color			
		dark	pastel	white	total
item_name	skirt	8	35	10	53
	dress	20	10	5	35
	shirt	14	7	28	49
	pants	20	2	5	27
	total	62	54	48	164

Figure 5.17 Cross tabulation of *sales* by *item.name* and *color*.

## 2. DATA CUBE

- All combinations with summarisations
- Here: measure attribute = quantity

		2	5	3	1	11		
		4	7	6	12	29		
		2	8	5	7	22		
color	dark	8	20	14	20	62	16	
	pastel	35	10	7	2	54	4	18
	white	10	5	28	5	48	21	45
	all	53	35	49	27	164	77	42
		skirt	dress	shirt	pants	all	all	clothes_size

Figure 5.18 Three-dimensional data cube.

- Number of ways to group tuples for aggregation  
 $3 \times 3 \times 4 = 36$
- With all  
 $4 \times 4 \times 5 = 80$
- For n dimensions, all value combinations on every group in the power set  $2^n$  are computed for  
 $2^3 = 8$  groups  
group by on each group

### 3. PIVOT

- Cross-tab is 2D view
- change dimensions used in cross-tab

### 4. SLICING

- Slice of data cube

### 5. ROLL UP

- Fine granularity → coarse granularity

### 6. DRILL DOWN

- Coarse granularity → fine granularity
- Must be generated

## Implementation of OLAP

- MOLAP : Multidimensional attr
- ROLAP : Relational
- HOLAP : Hybrid

```
select item_name, color, clothes_size, sum(quantity)
from sales
group by rollup(item_name, color, clothes_size);
```