

BIG DATA

UNIT - 5

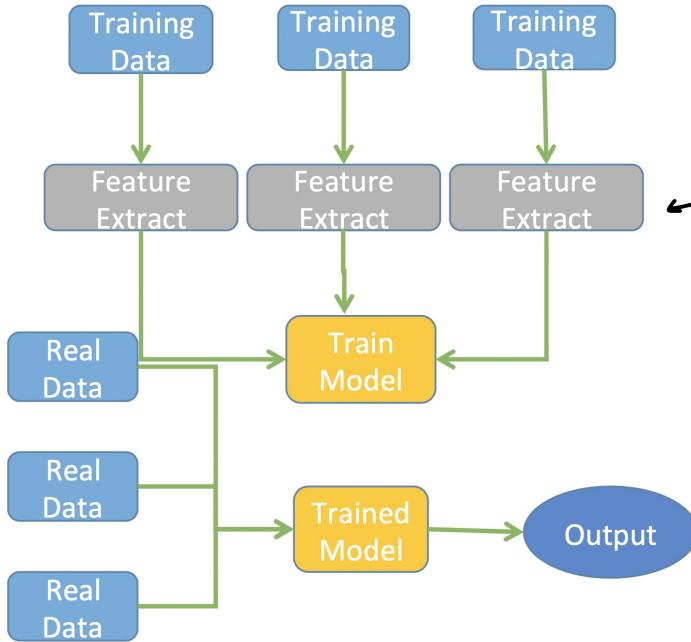
Advanced Analytics on
Big Data

feedback/corrections: vibha@pesu.pes.edu

VIBHA MASTI

Machine Learning

- Complex programs



Q: You want to write a program to separate out the rocks into different categories.

What will your approach be?



- Record colour, size, weight, thickness etc
- Cluster (unsupervised)
- Classify some training data into groups
- KNN / neural nets (supervised)

Supervised vs Unsupervised Learning

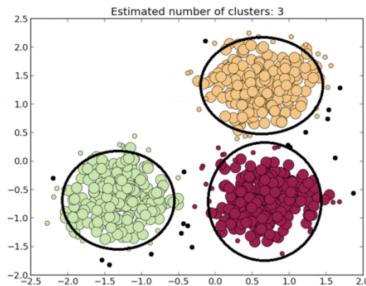
- Eg: grouping IPL batsmen
 - manually (opening, attacking): supervised
 - automatically: unsupervised

Q: Classify as supervised or unsupervised

- In Google News, grouping together similar articles.
unsupervised
- Determining if a particular credit card transaction is fraudulent
supervised
- Analyzing an image to determine if a lump is cancerous
supervised
- Recommending a product based on what the user buys
unsupervised
- Market segmentation: dividing customers into various groups
 - unsupervised - no pre-defined groups
 - supervised - pre-defined groups

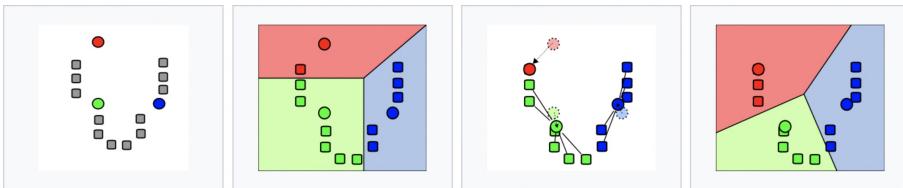
1. Scalable ML - K-means Clustering

- Partition dataset into k clusters
- Each training example is a point in d -dimensional space (d parameters / attributes)
- Use distance between points
 - Euclidean
 - Manhattan
 - Cosine
- Eg: 3-means



- Demonstration of k-means algorithm (iterative algorithm)

Demonstration of the standard algorithm



1. k initial "means" (in this case $k=3$) are randomly generated within the data domain (shown in color).

2. k clusters are created by associating every observation with the nearest mean. The partitions here represent the Voronoi diagram generated by the means.

3. The centroid of each of the k clusters becomes the new mean.

4. Steps 2 and 3 are repeated until convergence has been reached.

Algorithm

1. Select k random centres
2. Assign each data point to its closest cluster centre — form k clusters
3. Re-compute cluster centroids (average of each dimension)
4. If new centres different from old centres, back to step 2

- Q: • How can the k-means algorithm be modified to run with MapReduce?
• What is the output of Map and Reduce stages?

Hints:

- Iterative algorithm like page rank
- Which steps can be done in Map and which in Reduce?

Input: dataset - set of points - large
 k centroids - small

Map: read both input files
assign points to centroids

Output: <centroid, point>

key: centroid

Reduce: gets points associated with cluster
compute new centroids
Output: <new-centroid>

Loop: compare old and new centroids
check max iterations

Q: Given the following points

- 20, 30, 99, 102,
- 53, 9, 11, 54

Partition them into two clusters using k-means assuming initial centroids are 20, 30.

- Assume that each row of numbers is on a different machine
- Show what the keys and values are for one iteration of k-means

Machine 1

points: 20, 30, 99, 102
clusters: 20, 30

Machine 2

points: 53, 9, 11, 54
clusters: 20, 30

Mapper 1 output

<1, 20>
<2, 30>
<2, 99>
<2, 102>

Mapper 2 output

<2, 53>
<1, 9>
<1, 11>
<2, 54>

Reducer input

<1, [20, 9, 11]>
<2, [30, 99, 102, 53, 54]>

Reducer output

<1, 13.33> OR 13.33
<2, 67.6> OR 67.6

K-means Optimisations

1. Use of Combiners

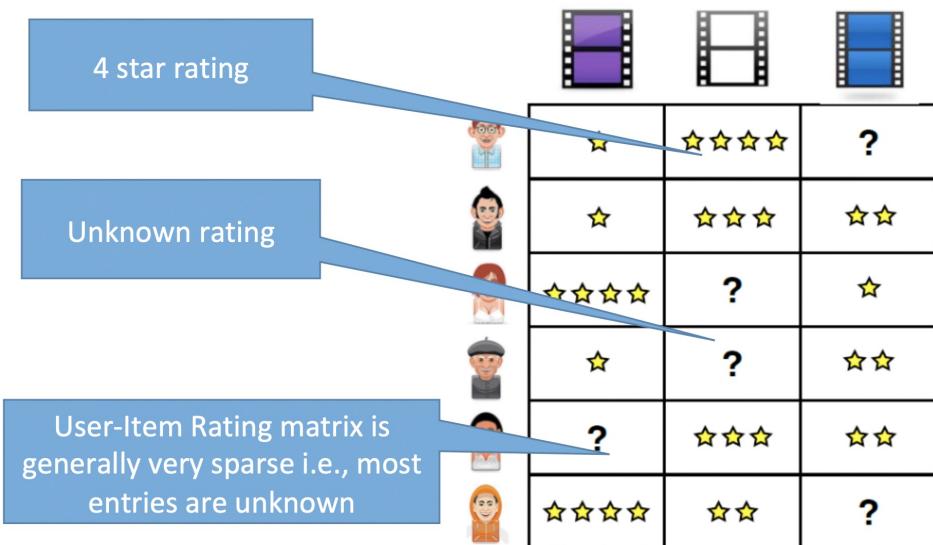
- compute local sums and count of the assigned clusters
- Sends `<centroid, <partial-sum, count>>`

2. Use of single Reducer

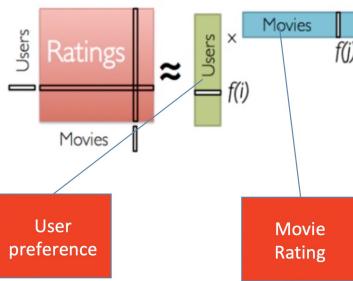
- Amount of data to reducers small
- Single reducer to tell if clusters changed
- Single output file

2. Scalable ML - Alternating Least Squares

- Collaborative filtering
- User-item relationships (rating of each item by each user)
- Sparse matrix



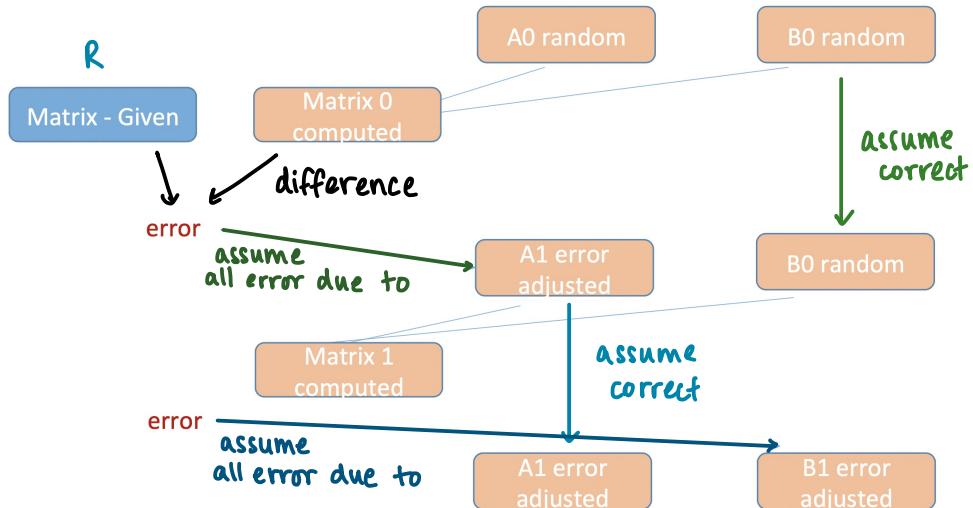
- User-Item rating matrix $R_{n \times m}$
- Try to write as a product of
 - User vector $A_{n \times 1}$
 - Item vector $B_{m \times 1}$
- Calculate A, B such that $R \approx AB$
- R : known
- A : unknown
- B : unknown
- r_{ij} : rating by user i for item j



- Approaches to factorisation
 - Gradient descent for optimisation problem (too slow)
 - Factorisation using alternating least squares
- Iterative algorithm with random initial assignments of A, B
- On i^{th} iteration, we have A_{i-1} and B_{i-1}
 - Assume B_{i-1} is correct and use to calculate best A_i
 - Assume A_i is correct and use to calculate best B_i
- Loop until convergence

ALS : i th Iteration

- Assume B_{i-1} correct
- Consider $R - A_i B_{i-1}^T$: error term
- $\|R - A_i B_{i-1}^T\|$ taken as $R - A_i B_{i-1}^T$ is a matrix and we want a number (absolute determinant)
- Find A_i that will minimise $\|R - A_i B_{i-1}^T\|$
- Shown that $A_i = (B_{i-1}^T B_{i-1})^{-1} B_{i-1}^T R^T$ (least squares regression estimate)
- Compute B_i from A_i



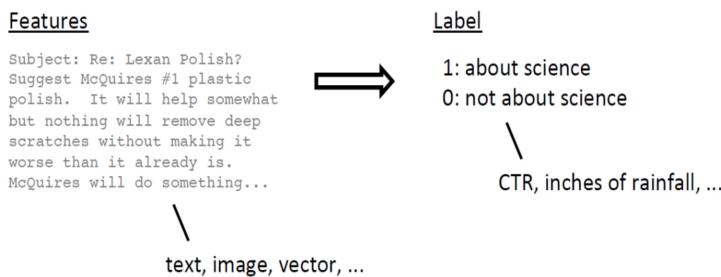
Q: ALS with Map-Reduce

- Use matrix multiplication with Map-Reduce to compute
 $A_i = (B_{i-1}^T B_{i-1})^{-1} B_{i-1}^T R^T$
- Similar for matrix inversion
- Compute-heavy

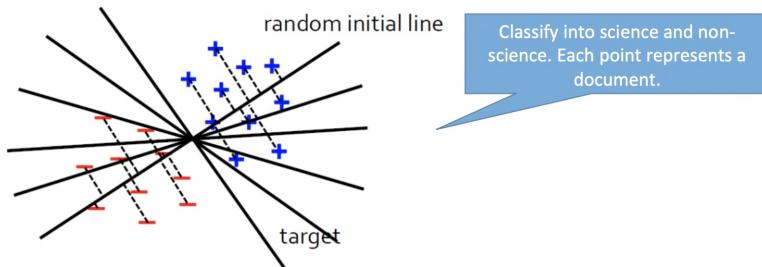
Spark MLLib

Text classification

- Supervised - given document, predict its topic

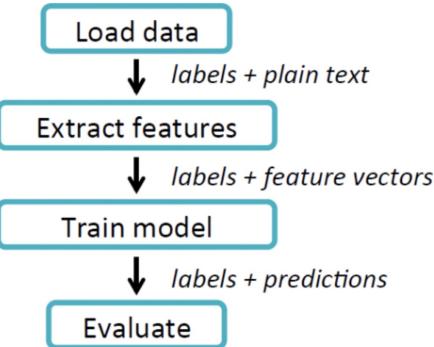


- Logistic regression: find separation between 2 classes
Support vector classifier: find line separating 2 classes

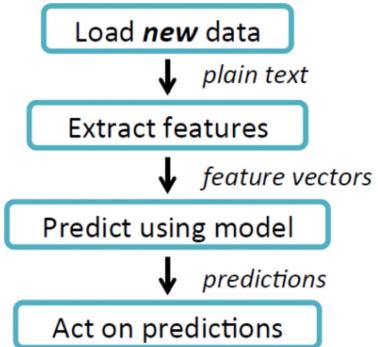


ML Workflow

TRAINING



TESTING/PRODUCTION



Pain Points

- Handle many RDDs (each feature, labels)
- ML pipeline must be written as a script (not modular)
- Train many models on diff splits with diff hyperparameters

Solve Challenges

- Make RDDs easier to read
 - break up the fields
 - Developers: program to extract features
- Read RDDs into dataframes
 - ML pipeline: transformers, estimators, evaluators
 - Parameter tuning: API

Dataframes

- KBs to PBs
- Wide variety of formats (Hive, existing RDDs)
- SOTA optimisation and code generation via Spark SQL catalyst optimiser
- APIs in Python, Java
- RDD + Schema + DSL (domain specific language)

Dataframe : RDD + Schema + DSL

Named columns with types

```
label: Double  
text: String  
words: Seq[String]  
features: Vector  
prediction: Double
```

label	text	words	features
0	This is ...	["This", "is", ...]	[0.5, 1.2, ...]
0	When we ...	["When", "..."]	[1.9, -0.8, ...]
1	Knuth was ...	["Knuth", "..."]	[0.0, 8.7, ...]
0	Or you ...	["Or", "you", ...]	[0.1, -0.6, ...]

Domain-Specific Language

```
# Select science articles  
sciDocs =  
    data.filter("label" == 1)  
  
# Scale labels  
data("label") * 0.5
```

Spark ML Pipelines

- Automates process of defining models

ML Pipeline

1. Transformers

- Extract features from dataframe
- Features stored in new dataframe

2. Estimators

- ML algorithms (standard defined or user defined)

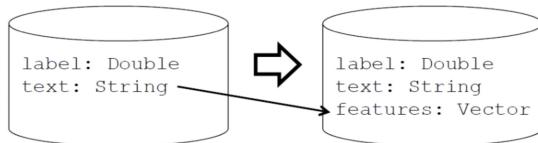
3. Evaluators

- Compute predictions
- Estimate error metrics
- Tune algorithm parameters
- Evaluator depends on estimator (logistic regression, decision trees)

1. Transformers

- Extract features
- Input: dataframe, Output: dataframe

```
def transform(DataFrame): DataFrame
```



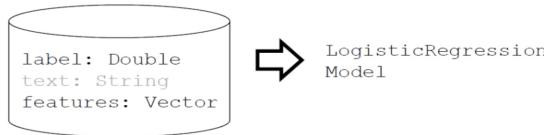
Label	Text
0	
1	

Label	Text	Features
0		
1		

2. Estimator

- Train model

```
def fit(DataFrame): Model
```

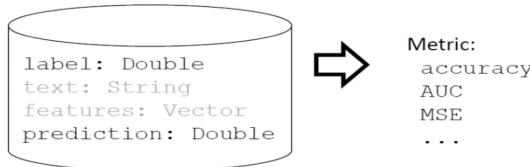


Label	Text	Features
0		
1		

3. Evaluator

- Metrics

```
def evaluate(DataFrame): Double
```

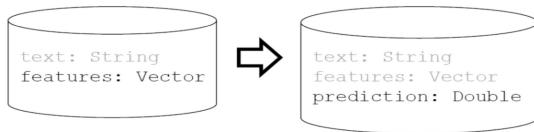


Label	Text	Features	Prediction
0			
1			

2. Predict Using Model (Test Phase)

Model is a type of Transformer

```
def transform(DataFrame): DataFrame
```

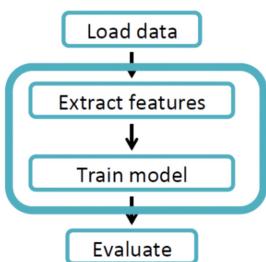


Text	Features

Text	Features	Prediction

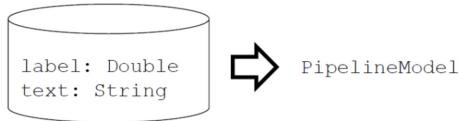
TRAINING PIPELINE

TRAINING



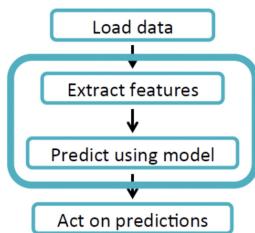
Pipeline is a type of Estimator

```
def fit(DataFrame): Model
```



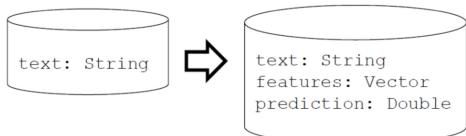
TESTING PIPELINE

TESTING/PRODUCTION

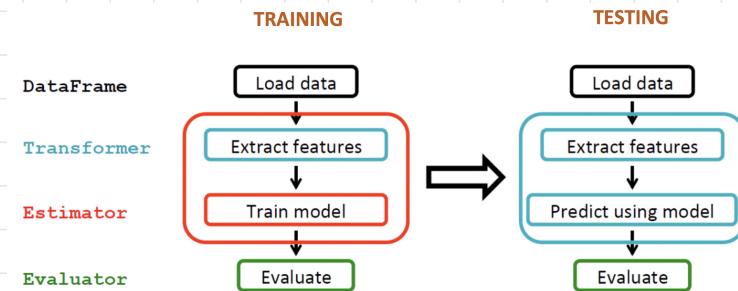


PipelineModel is a type of Transformer

```
def transform(DataFrame): DataFrame
```

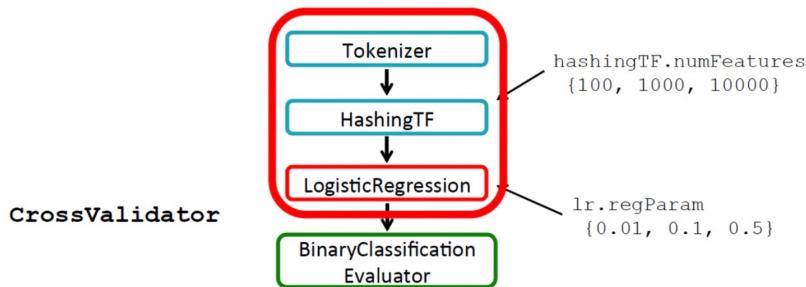


All Together



Parameter Tuning

- Given: estimator, parameter grid, evaluator
 - find best parameters



- Q: Suppose we have a dataset in which each line has a recording of a noise, and its classification
- E.g., `<bell.wav>`, bell
- What would be the input DataFrame be?

- Suppose we want to recognize sounds by
- Extracting the frequencies from the wav file
 - Gaussian model
 - Find the average frequency of each sound
 - For a new sound, calculate average frequency
 - Find closest matching sound

- What are the DataFrames, Evaluators, etc needed?

1. Input dataframe

`(bell.wav, bell)`

2. Feature dataframe

`(bell.wav, bell, frequencies)`

3. Transformer

`<bell.wav, bell, avg-frequency>`

4. Model

- train (feature-dataframe) : find avg freq for each sound type
- predict (predict-dataframe) : find closest matching sound

deep learning

- Iterative - matrix multiplication
- Challenges for DL
 - data on HDFS, train on Spark

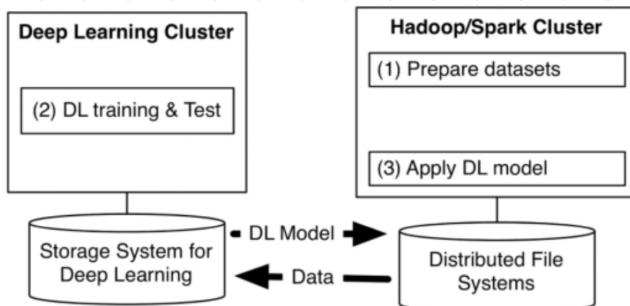


Figure 1: ML Pipeline with multiple programs on separated clusters

Tensorflow on Spark

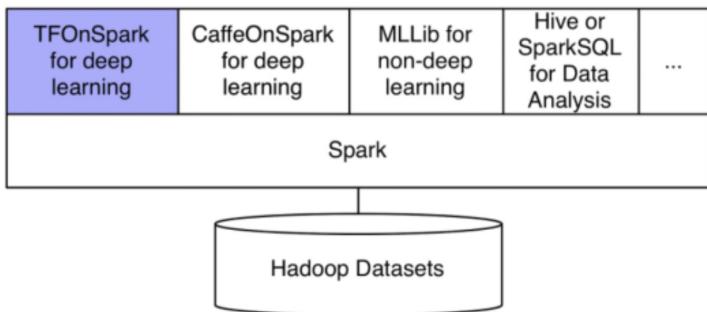


Figure 2: TensorFlowOnSpark for deep learning on Spark clusters

- Supports model & data parallelism
- TF runs on Spark executors
- Read data directly from HDFS
- RDMA: fast network transfers

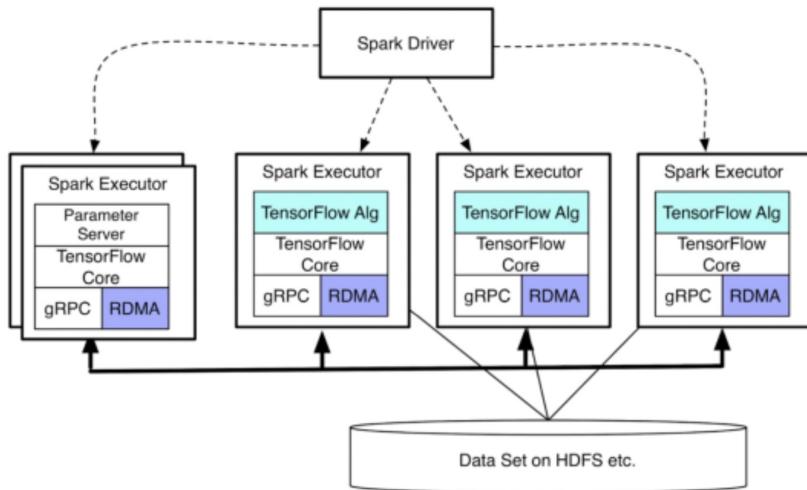


Figure 3: TensorFlowOnSpark system architecture