



# OPERATING SYSTEMS

## CPU Scheduling

---

**Chandravva Hebbi**

Department of Computer Science

- The slides/diagrams in this course are an **adaptation**, **combination**, and **enhancement** of material from the following resources and persons:
  1. Slides of Operating System Concepts, Abraham Silberschatz, Peter Baer Galvin, Greg Gagne - 9<sup>th</sup> edition 2013 and some slides from 10<sup>th</sup> edition 2018
  2. Some conceptual text and diagram from Operating Systems - Internals and Design Principles, William Stallings, 9<sup>th</sup> edition 2018
  3. Some presentation transcripts from A. Frank – P. Weisberg
  4. Some conceptual text from Operating Systems: Three Easy Pieces, Remzi Arpaci-Dusseau, Andrea Arpaci Dusseau

# OPERATING SYSTEMS

---

## Process Scheduling

**Chandravva Hebbi**

Department of Computer Science

# OPERATING SYSTEMS

## Basic Concepts

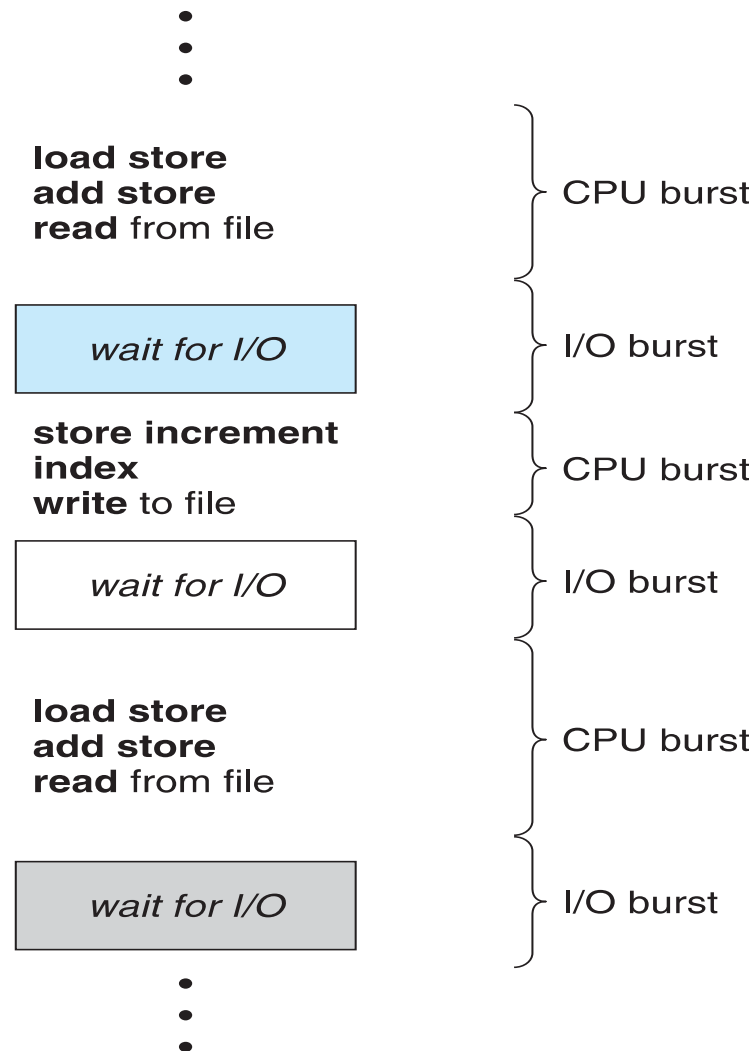
---




- ❑ In a system with a single CPU core, only one process can run at a time. Others must wait until the CPU is free and can be rescheduled.
- ❑ The objective of multiprogramming is to have some process running at all times, to maximize CPU utilization.
- ❑ Several processes are kept in memory at one time.
- ❑ When one process has to wait, the operating system takes the CPU away from that process and gives the CPU to another process. This pattern continues.
- ❑ Every time one process has to wait, another process can take over use of the CPU. On a multicore system, this concept of keeping the CPU busy is extended to all processing cores on the system.

- ❑ The idea is relatively simple. A process is executed until it must wait, typically for the completion of some I/O request.
- ❑ In a simple computer system, the CPU then just sits idle. All this waiting time is wasted; no useful work is accomplished. With multiprogramming, we try to use this time productively.
- ❑ Scheduling of this kind is a fundamental operating-system function.
- ❑ Almost all computer resources are scheduled before use. The CPU is, of course, one of the primary computer resources. Thus, its scheduling is central to operating-system design.

- ❑ Maximum CPU utilization obtained with multiprogramming
- ❑ CPU–I/O Burst Cycle – Process execution consists of a **cycle** of CPU execution and I/O wait
- ❑ **CPU burst** followed by **I/O burst**
- ❑ CPU burst distribution is of main concern



- The durations of CPU bursts have been measured extensively. Although they vary greatly from process to process and from computer to computer, they tend to have a frequency curve similar to that shown in the Figure.
- An I/O-bound program typically has many short CPU bursts. A CPU-bound program might have a few long CPU bursts. This distribution can be important when implementing a CPU-scheduling algorithm.



- ❑ **Short-term scheduler** selects from among the processes in ready queue, and allocates the CPU to one of them
  - ❑ Queue may be ordered in various ways
- ❑ CPU scheduling decisions may take place when a process:
  1. Switches from running to waiting state
  2. Switches from running to ready state
  3. Switches from waiting to ready
  4. Terminates
- ❑ Scheduling under 1 and 4 is **nonpreemptive**
- ❑ All other scheduling is **preemptive**
  - ❑ Consider access to shared data
  - ❑ Consider preemption while in kernel mode
  - ❑ Consider interrupts occurring during crucial OS activities



- Under non-preemptive scheduling, once the CPU has been allocated to a process, the process keeps the CPU until it releases it either by terminating or by switching to the waiting state.
- Virtually all modern operating systems including Windows, macOS, Linux, and UNIX use pre-emptive scheduling algorithms.

- Unfortunately, pre-emptive scheduling can result in race conditions when data are shared among several processes. **Ex:** While one process is updating the shared data, it is pre-empted so that the second process can run. The second process then tries to read the data, which are in an inconsistent state.
- A pre-emptive kernel requires mechanisms such as mutex locks to prevent race conditions when accessing shared kernel data structures. **Most modern operating systems are now fully pre-emptive when running in kernel mode.**

- ❑ Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
  - ❑ switching context
  - ❑ switching to user mode
  - ❑ jumping to the proper location in the user program to restart that program
- ❑ **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running

- ❑ **CPU utilization** – keep the CPU as busy as possible
- ❑ **Throughput** – # of processes that complete their execution per time unit
- ❑ **Turnaround time** – amount of time to execute a particular process (performance metric)
- ❑ **Waiting time** – amount of time a process has been waiting in the ready queue
- ❑ **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

# OPERATING SYSTEMS

## Scheduling Algorithm Optimization Criteria

---

- ❑ Max CPU utilization
- ❑ Max throughput
- ❑ Min turnaround time
- ❑ Min waiting time
- ❑ Min response time

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Suppose that the processes arrive in the order:  $P_1, P_2, P_3$   
The Gantt Chart for the schedule is:



<sup>0</sup>Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$

- Average waiting time:  $(0 + 24 + 27)/3 = 17$

- Suppose that the processes arrive in the order:  $P_2, P_3, P_1$   
The Gantt Chart for the schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- **Convoy effect** - short process behind long process
  - Consider one CPU-bound and many I/O-bound processes

# OPERATING SYSTEMS

## Take Home Assignment



Consider the following set of processes, with the length of the CPU burst given in milliseconds

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	10	3
$P_2$	1	1
$P_3$	2	3
$P_4$	1	4
$P_5$	5	2

The processes are assumed to have arrived in the order  $P_1, P_2, P_3, P_4, P_5$ , all at time 0.

- Draw four Gantt charts that illustrate the execution of these processes using the following scheduling algorithms: FCFS, SJF Scheduling algorithm
- What is the turnaround time of each process for each of the scheduling algorithms
- What is the waiting time of each process for each of these scheduling algorithms?



# OPERATING SYSTEMS

## Take Home Assignment

---

Which of the following scheduling algorithms could result in starvation?

- a. First-come, first-served
- b. Shortest job first
- c. Round robin
- d. Priority





**THANK YOU**

---

**Chandravva Hebbi**

Department of Computer Science Engineering

**[chandravvahebbi@pes.edu](mailto:chandravvahebbi@pes.edu)**