



OPERATING SYSTEMS

Multi-Processor Scheduling and Real-Time CPU Scheduling

Chandravva Hebbi

Department of Computer Science

OPERATING SYSTEMS

Slides Credits for all PPTs of this course



- The slides/diagrams in this course are an **adaptation**, **combination**, and **enhancement** of material from the following resources and persons:
 1. Slides of Operating System Concepts, Abraham Silberschatz, Peter Baer Galvin, Greg Gagne - 9th edition 2013 and some slides from 10th edition 2018
 2. Some conceptual text and diagram from Operating Systems - Internals and Design Principles, William Stallings, 9th edition 2018
 3. Some presentation transcripts from A. Frank – P. Weisberg
 4. Some conceptual text from Operating Systems: Three Easy Pieces, Remzi Arpaci-Dusseau, Andrea Arpaci Dusseau

OPERATING SYSTEMS

Multi-Processor Scheduling Real-Time CPU Scheduling

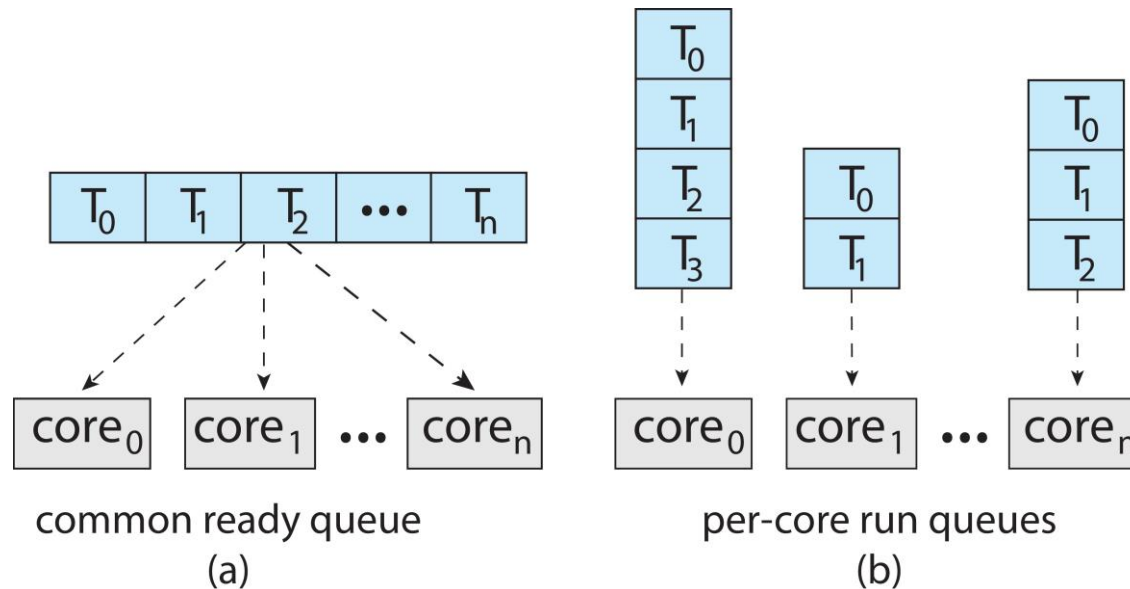
Chandravva Hebbi

Department of Computer Science

- ❑ CPU scheduling more complex when multiple CPUs are available, load sharing will be possible
- ❑ **Asymmetric multiprocessing** – only one processor accesses the system data structures, alleviating the need for data sharing
- ❑ **Symmetric multiprocessing (SMP)** – each processor is self-scheduling, all processes in common ready queue, or each has its own private queue of ready processes
 - ❑ Currently, most common.
- ❑ Virtually all modern operating systems support SMP, including Windows, Linux, and Mac OS X.

- ❑ what happens to cache memory when a process has been running on a specific processor?
 - ❑ Populated its buffer with data
 - ❑ What happens if process migrates?
- ❑ **Processor affinity** – process has affinity for processor on which it is currently running
 - ❑ **soft affinity:**
 - ❑ OS keeps the process with processor, but not guaranteed.
 - ❑ **hard affinity**
 - ❑ OS does not allow process to migrate between processors.
 - ❑ Linux implements soft affinity
 - ❑ The **sched_setaffinity()** system call, which supports hard affinity
 - ❑ Variations including **processor sets**


- Symmetric multiprocessing (SMP) is where each processor is self scheduling.
- All threads may be in a common ready queue (a)
- Each processor may have its own private queue of threads (b)



OPERATING SYSTEMS

NUMA and CPU Scheduling

- The main-memory architecture of a system can affect processor affinity issues.



If the operating system's CPU scheduler and memory-placement algorithms work together, then a process that is assigned affinity to a particular CPU can be allocated memory on the board where that CPU resides

- ❑ If SMP, need to keep all CPUs loaded for efficiency.
- ❑ Load balancing required when processors have their own queue

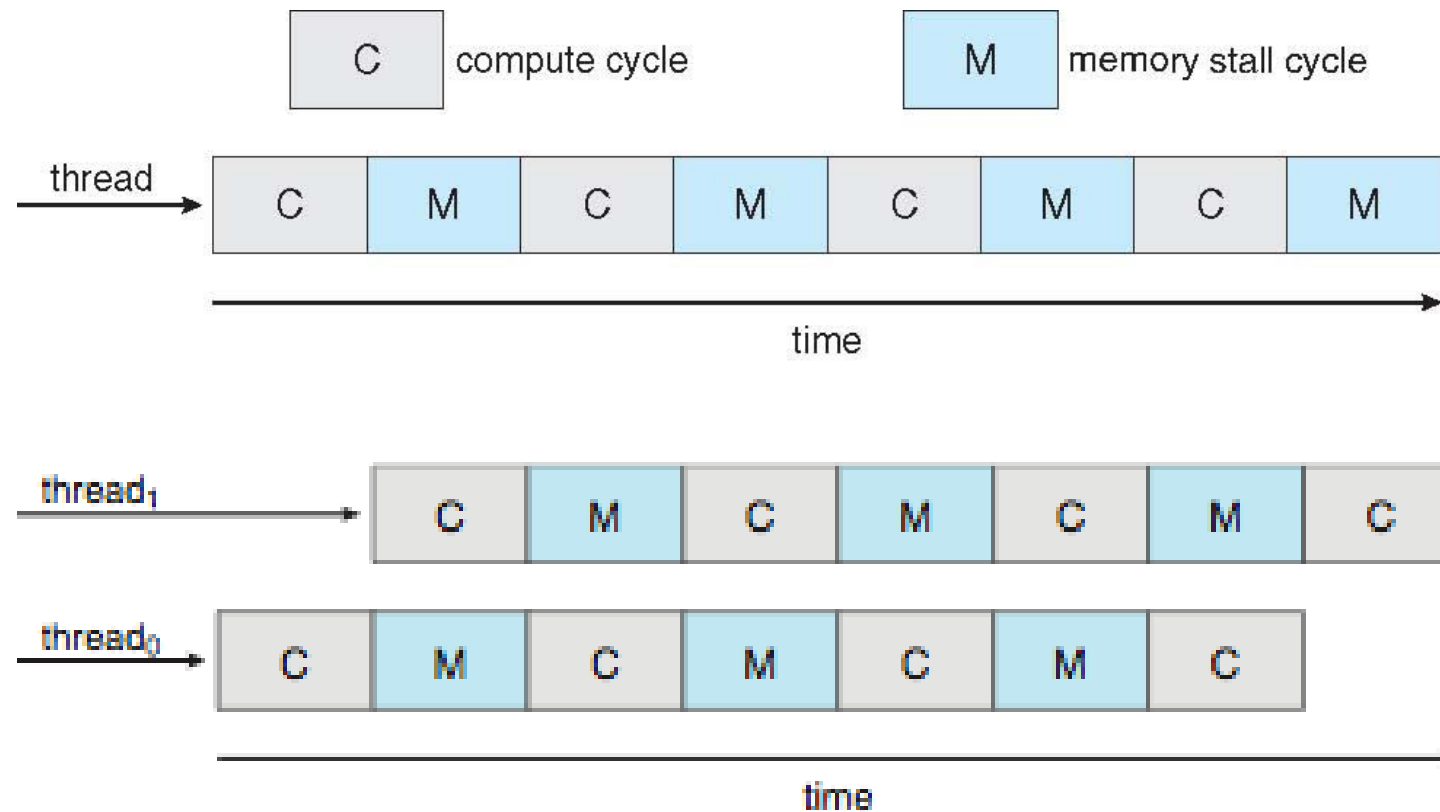
Approaches to load balancing

- ❑ **Load balancing** attempts to keep workload evenly distributed
- ❑ **Push migration** – periodic task checks load on each processor, and if found pushes task from overloaded CPU to other CPUs
- ❑ **Pull migration** – idle processors pulls waiting task from busy processor
- ❑ Load balancing often counteracts the benefits of process affinity

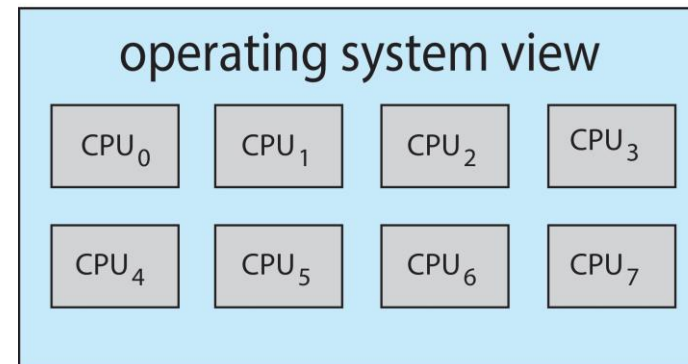
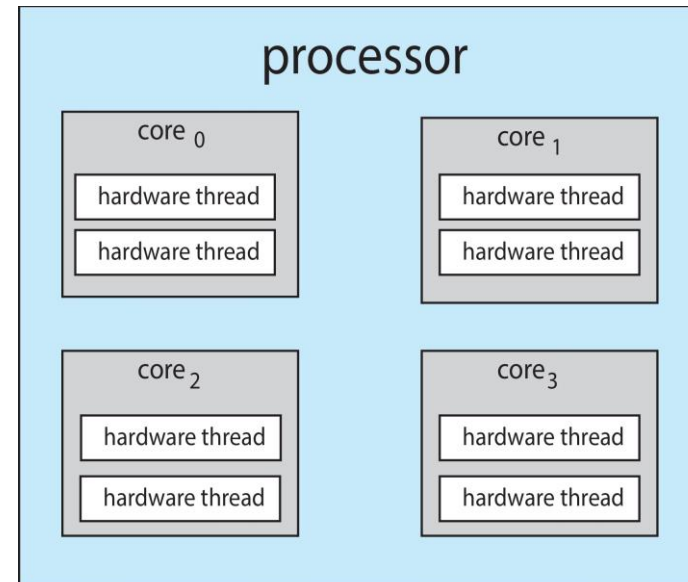
- Recent trend to place multiple processor cores on same physical chip
- Faster and consumes less power
- Multiple threads per core also growing
 - Takes advantage of memory stall to make progress on another thread while memory retrieve happens
- Memory stall is the situation when a processor accesses memory, it spends a significant amount of time waiting for the data to become available.
- Each core has > 1 hardware threads. If one thread has a memory stall, switch to another thread!

OPERATING SYSTEMS

Multithreaded Multicore System



- **Chip-multithreading (CMT)** assigns each core multiple hardware threads. (Intel refers to this as **hyperthreading**.)
- On a quad-core system with 2 hardware threads per core, the operating system sees 8 logical processors.



Two ways to multithread a processing core:

- **Coarse grained** multithreading

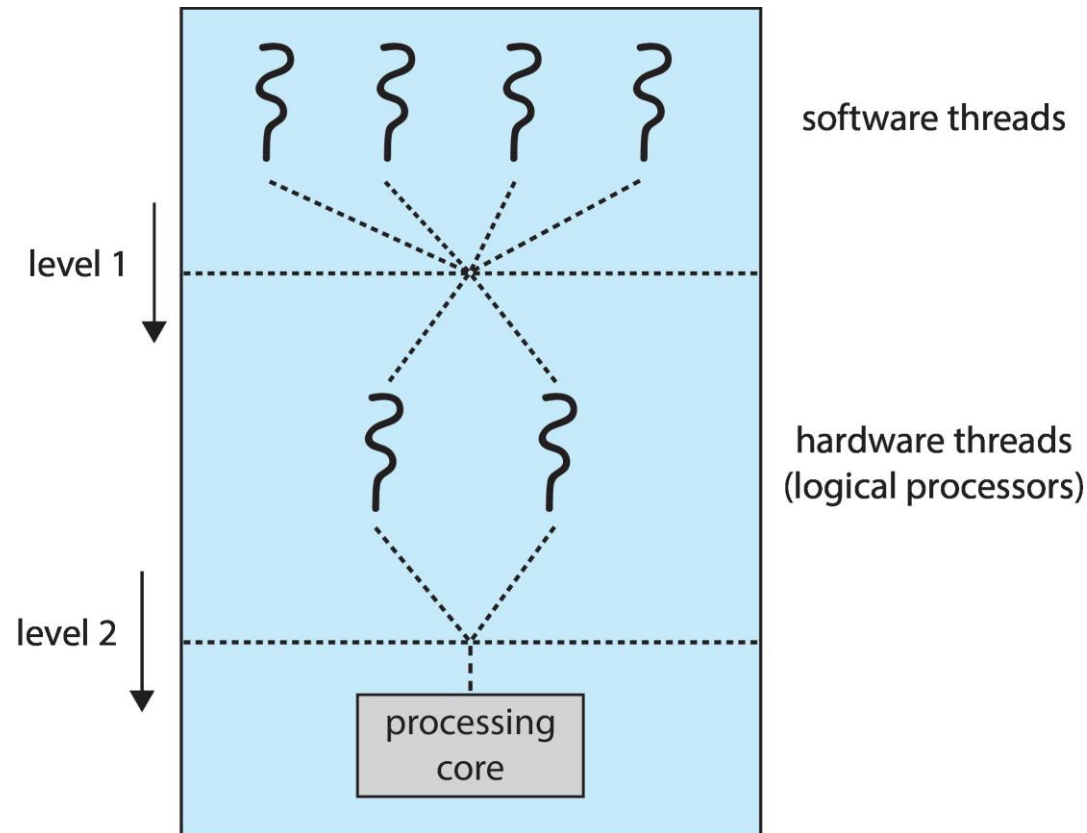
- A thread executes on a processor until a long-latency event such as a memory stall occurs.
- Because of the long latency processor switches to another thread.
- Cost of switching is high. Why?

- **fine-grained** multithreading

- Switching between the threads at finer level of granularity
- Cost of switching low.

Multithreaded multicore processor needs two levels of scheduling:

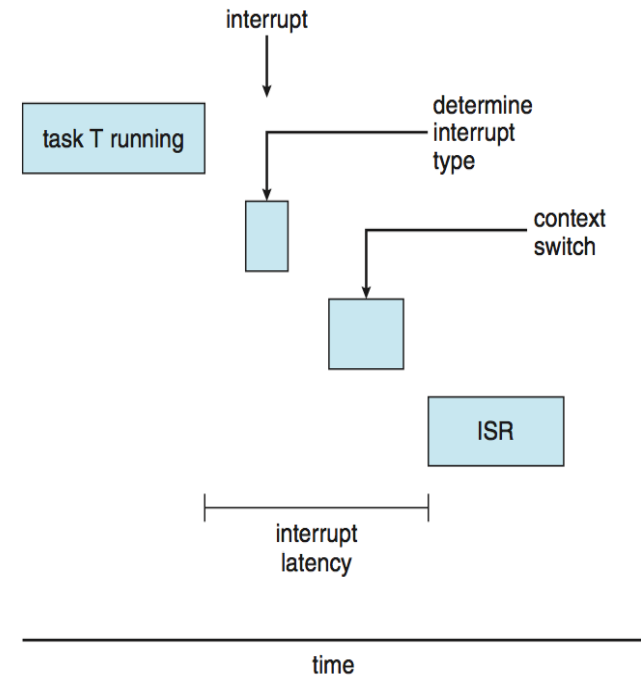
1. The operating system deciding which software thread to run on a logical CPU
2. How each core decides which hardware thread to run on the physical core.



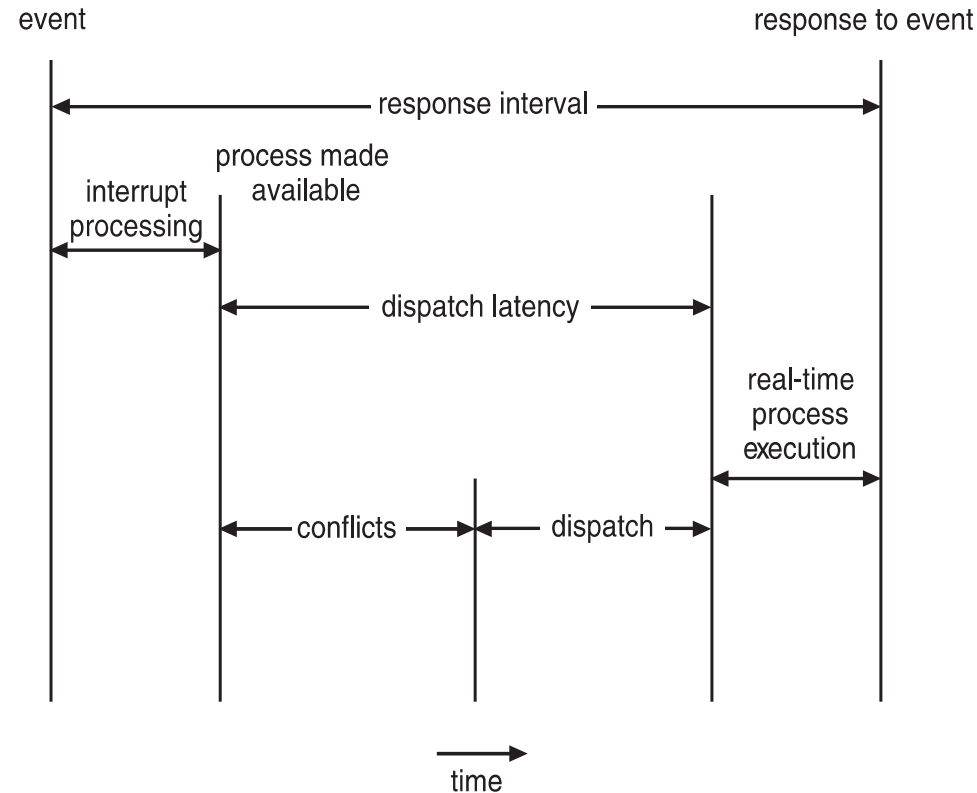
OPERATING SYSTEMS

Real-Time CPU Scheduling


- ❑ Can present obvious challenges
- ❑ **Soft real-time systems** – no guarantee as to when critical real-time process will be scheduled
- ❑ **Hard real-time systems** – task must be serviced by its deadline.
- ❑ Different events have different latencies.
- ❑ Two types of latencies affect performance
 1. Interrupt latency – time from arrival of interrupt to start of routine that services interrupt
 2. Dispatch latency – time for scheduler to take current process off CPU and switch to another



- Conflict phase of dispatch latency:
1. Preemption of any process running in kernel mode
 2. Release by low-priority process of resources needed by high-priority processes



- ❑ For real-time scheduling, scheduler must support preemptive, priority-based scheduling
 - ❑ But only guarantees soft real-time
- ❑ For hard real-time must also provide ability to meet deadlines
- ❑ Processes have new characteristics: **periodic**-ones require CPU at constant intervals
 - ❑ Has processing time t , deadline d , period p
 - ❑ $0 \leq t \leq d \leq p$
 - ❑ **Rate** of periodic task is $1/p$
- ❑ Schedulers can take advantage of these characteristics and assign priorities according to a process's deadline or rate requirements.



OPERATING SYSTEMS

Rate Monotonic Scheduling

- ❑ Priority based algorithm that belongs to the static priority scheduling category for Real Time Operating Systems.
- ❑ A priority is assigned based on the inverse of its period
- ❑ Shorter periods = higher priority;
- ❑ Longer periods = lower priority


Example

P1 and P2 are 50 and 100, respectively—that is, $p_1 = 50$ and $p_2 = 100$.

The processing times are $t_1 = 20$ for P1 and $t_2 = 35$ for P2. The deadline for each process requires that it complete its CPU burst by the start of its next period.

OPERATING SYSTEMS

Missed Deadlines with Rate Monotonic Scheduling




- CPU utilization- t_i/p_i
- $P_1=20/50=.40$ $P_2=35/100=.35$ total utilization=.75
which is 75%

Case 1

- P_2 is given higher priority than P_1 .
- P_1 completes at 55, it misses the deadline

Case 2


- P_1 is given highest priority as P_1 time is smaller than P_2 .
- P_1 and P_2 complete by 75.
- CPU is idle till 100



OPERATING SYSTEMS

Missed Deadlines with Rate Monotonic Scheduling

Consider Processes P1 and P2 with P1=50, t1=25 and P2=80, t2=35




Worst case CPU utilization for N processes is: $N(2^{1/N}-1)$

CPU utilization falls as the number of process approaches to infinity

- Priorities are assigned dynamically according to deadlines:
 - the earlier the deadline, the higher the priority;
 - the later the deadline, the lower the priority
- Process must announce its deadline when it becomes runnable

Example

- Consider Processes P1 and P2 with $P1=50$, $t1=25$ and $P2=80$, $t2=35$



- T shares are allocated among all processes in the system
- An application receives N shares where $N < T$
- This ensures each application will receive N / T of the total processor time.

As an example,

assume that a total of $T = 100$ shares is to be divided among three processes A , B , and C .

A is assigned 50 shares, B is assigned 15 shares, and C is assigned 20 shares.

What happens if new process D request for 30 shares

- ❑ The POSIX.1b standard
- ❑ API provides functions for managing real-time threads
- ❑ Defines two scheduling classes for real-time threads:
 1. SCHED_FIFO - threads are scheduled using a FCFS strategy with a FIFO queue. There is no time-slicing for threads of equal priority
 2. SCHED_RR - similar to SCHED_FIFO except time-slicing occurs for threads of equal priority
- ❑ Defines two functions for getting and setting scheduling policy:
 1. `pthread_attr_getsched_policy(pthread_attr_t *attr, int *policy)`
 2. `pthread_attr_setsched_policy(pthread_attr_t *attr, int policy)`



THANK YOU

Chandravva Hebbi

Department of Computer Science Engineering

chandravvahebbi @pes.edu