



# OPERATING SYSTEMS

## IPC— Shared Memory & Message Passing, Pipes – Named and Unnamed

---

**Chandravva Hebbi**

Department of Computer Science

- The slides/diagrams in this course are an **adaptation**, **combination**, and **enhancement** of material from the following resources and persons:
  1. Slides of Operating System Concepts, Abraham Silberschatz, Peter Baer Galvin, Greg Gagne - 9<sup>th</sup> edition 2013 and some slides from 10<sup>th</sup> edition 2018
  2. Some conceptual text and diagram from Operating Systems - Internals and Design Principles, William Stallings, 9<sup>th</sup> edition 2018
  3. Some presentation transcripts from A. Frank – P. Weisberg
  4. Some conceptual text from Operating Systems: Three Easy Pieces, Remzi Arpaci-Dusseau, Andrea Arpaci Dusseau

# OPERATING SYSTEMS

---

## IPC— Shared Memory & Message Passing, Pipes – Named and Unnamed

**Chandravva Hebbi**

Department of Computer Science

# OPERATING SYSTEMS

## Interprocess Communication

---



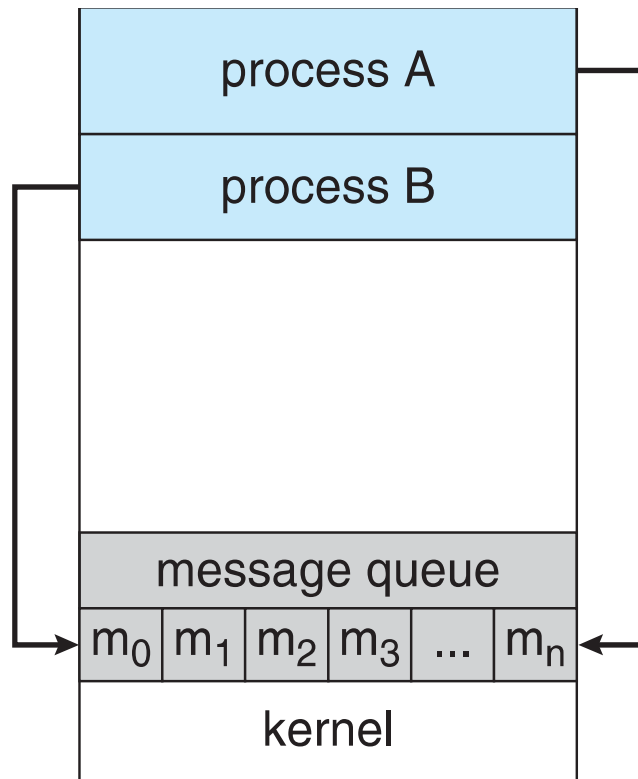
- ❑ Processes within a system may be *independent* or *cooperating*
- ❑ Cooperating process can affect or be affected by other processes, including sharing data
- ❑ Reasons for cooperating processes:
  - ❑ Information sharing
  - ❑ Computation speedup
  - ❑ Modularity
  - ❑ Convenience
- ❑ Cooperating processes need **interprocess communication (IPC)**
- ❑ Two models of IPC
  - ❑ **Shared memory**
  - ❑ **Message passing**

- ❑ **Independent** process cannot affect or be affected by the execution of another process
- ❑ **Cooperating** process can affect or be affected by the execution of another process
- ❑ Advantages of process cooperation
  - ❑ Information sharing
  - ❑ Computation speed-up
  - ❑ Modularity
  - ❑ Convenience


# OPERATING SYSTEMS

## Communication Models

- Two models of IPC
  - a) Message passing and
  - b) Shared memory



(a)



(b)

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
  - **unbounded-buffer** places no practical limit on the size of the buffer
    - Consumer may have to wait for new items, but the producer can always produce new items
  - **bounded-buffer** assumes that there is a fixed buffer size
    - Consumer must wait if the buffer is empty; producer must wait if the buffer is full

- ❑ Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    ...
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

- ❑ Shared buffer is implemented as a circular array with 2 logical pointers: **in** and **out**
- ❑ Buffer is empty when **in == out**; buffer is full when  $((\text{in} + 1) \% \text{BUFFER\_SIZE}) == \text{out}$
- ❑ Variable **in** points to the next free position in the buffer; **out** points to the first full position in the buffer
- ❑ Solution is correct, but can only use  $\text{BUFFER\_SIZE}-1$  elements



```
item next_produced;
while (true) {
    /* produce an item in next_produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

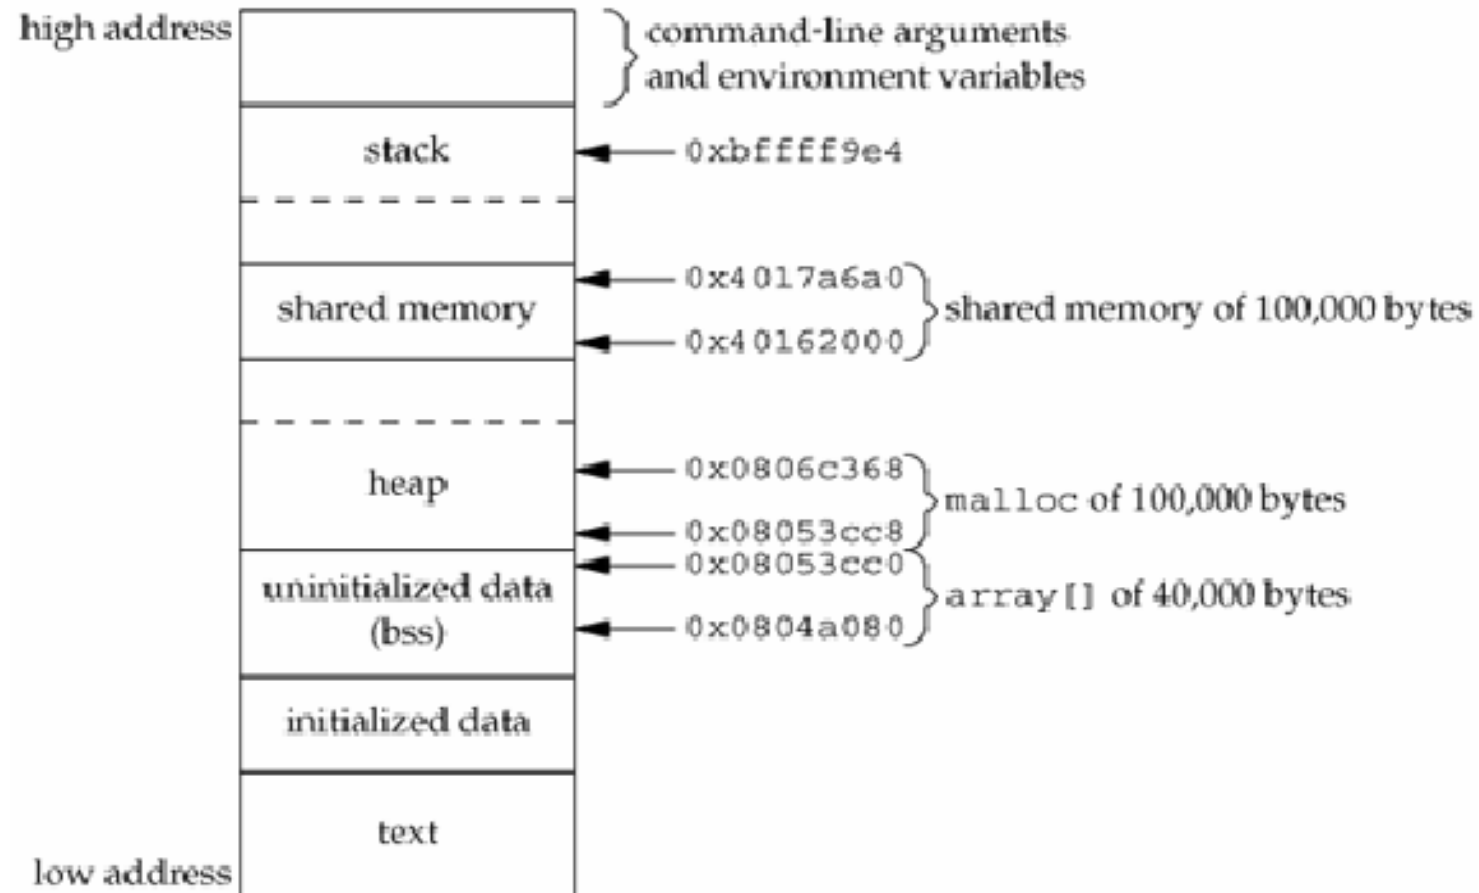
```
item next_consumed;
while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next_consumed */
}
```

- ❑ An area of memory shared among the processes that wish to communicate
- ❑ The communication is under the control of the users processes not the operating system.
- ❑ Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.

- ❑ Shared memory allows two or more processes to share a given region of memory.
- ❑ Shared memory is the fastest form of IPC, because the data does not need to be copied between the client and the server.
- ❑ The only trick in using shared memory is synchronizing access to a given region among multiple processes.
- ❑ If the server is placing data into a shared memory region, the client shouldn't try to access the data until the server is done.
- ❑ Often, semaphores are used to synchronize shared memory access. (record locking can also be used.)

### *Memory layout on an Intel-based Linux system*



- ❑ Mechanism for processes to communicate and to synchronize their actions
- ❑ Message system – processes communicate with each other without resorting to shared variables
- ❑ IPC facility provides two operations:
  - ❑ **send**(*message*)
  - ❑ **receive**(*message*)
- ❑ The *message* size is either fixed or variable

- ❑ If processes  $P$  and  $Q$  wish to communicate, they need to:
  - ❑ Establish a **communication link** between them
  - ❑ Exchange messages via send/receive
- ❑ Implementation issues:
  - ❑ How are links established?
  - ❑ Can a link be associated with more than two processes?
  - ❑ How many links can there be between every pair of communicating processes?
  - ❑ What is the capacity of a link?
  - ❑ Is the size of a message that the link can accommodate fixed or variable?
  - ❑ Is a link unidirectional or bi-directional?

- Implementation of communication link
  - Physical:
    - ▶ Shared memory
    - ▶ Hardware bus
    - ▶ Network
  - Logical:
    - ▶ Direct or indirect
    - ▶ Synchronous or asynchronous
    - ▶ Automatic or explicit buffering



- Processes must name each other explicitly:
  - **send** ( $P, message$ ) – send a message to process P
  - **receive**( $Q, message$ ) – receive a message from process Q
- Properties of communication link
  - Links are established automatically
  - A link is associated with exactly one pair of communicating processes
  - Between each pair there exists exactly one link
  - The link may be unidirectional, but is usually bi-directional

- ❑ Messages are directed and received from mailboxes (also referred to as ports)
  - ❑ Each mailbox has a unique id
  - ❑ Processes can communicate only if they share a mailbox
- ❑ Properties of communication link
  - ❑ Link established only if processes share a common mailbox
  - ❑ A link may be associated with many processes
  - ❑ Each pair of processes may share several communication links
  - ❑ Link may be unidirectional or bi-directional

- Operations
  - create a new mailbox (port)
  - send and receive messages through mailbox
  - destroy a mailbox
- Primitives are defined as:
  - send**(*A, message*) – send a message to mailbox A
  - receive**(*A, message*) – receive a message from mailbox A

- Mailbox sharing
  - $P_1$ ,  $P_2$ , and  $P_3$  share mailbox A
  - $P_1$  sends;  $P_2$  and  $P_3$  receive
  - Who gets the message?
- Solutions
  - Allow a link to be associated with at most two processes
  - Allow only one process at a time to execute a receive operation
  - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

- ❑ Message passing may be either blocking or non-blocking
- ❑ **Blocking** is considered **synchronous**
  - **Blocking send** -- the sender is blocked until the message is received
  - **Blocking receive** -- the receiver is blocked until a message is available
- ❑ **Non-blocking** is considered **asynchronous**
  - **Non-blocking send** -- the sender sends the message and continue
  - **Non-blocking receive** -- the receiver receives:
    - A valid message, or
    - Null message
- ❑ Different combinations possible
  - If both send and receive are blocking, we have a **rendezvous** between the sender and the receiver

### ❑ Producer-consumer becomes trivial

```
message next_produced;  
while (true) {  
    /* produce an item in next_produced */  
    send(next_produced);  
}
```


```
message next_consumed;  
while (true) {  
    receive(next_consumed);  
  
    /* consume the item in next_consumed */  
}
```

- Queue of messages attached to the link (direct or indirect); messages reside in a temporary queue.
- Queues can be implemented in one of three ways
  1. Zero capacity – no messages are queued on a link.  
Sender must wait for receiver (rendezvous)
  2. Bounded capacity – finite length of  $n$  messages  
Sender must wait if link full
  3. Unbounded capacity – infinite length  
Sender never waits
- Zero-capacity case is sometimes referred to as a message system with no buffering; other cases are referred to as systems with automatic buffering

- ❑ Acts as a conduit allowing two processes to communicate
- ❑ Issues:
  - ❑ Is communication unidirectional or bidirectional?
  - ❑ In the case of two-way communication, is it half or full-duplex?
  - ❑ Must there exist a relationship (i.e., ***parent-child***) between the communicating processes?
  - ❑ Can the pipes be used over a network?
- ❑ Ordinary (Unnamed) pipes – cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.
- ❑ Named pipes – can be accessed without a parent-child relationship.



- ❑ Ordinary Pipes allow communication in standard producer-consumer style
- ❑ Producer writes to one end (the **write-end** of the pipe)
- ❑ Consumer reads from the other end (the **read-end** of the pipe)
- ❑ Ordinary pipes are therefore unidirectional
- ❑ Require parent-child relationship between communicating processes



- ❑ Windows calls these **anonymous pipes**
- ❑ Used by the shell; not used very often by application programs
  - Main limitation is processes need to be related

- ❑ Named Pipes are more powerful than ordinary pipes
- ❑ Communication is bidirectional
- ❑ No parent-child relationship is necessary between the communicating processes
- ❑ Several processes can use the named pipe for communication
- ❑ Provided on both UNIX and Windows systems.
- ❑ Limitations of named pipes:
  - ❑ Two pipes needed for two way communication
  - ❑ FIFO: Data once retrieved is removed from pipe

- ❑ On UNIX Systems:
  - ❑ Named Pipes are referred to as FIFOs
  - ❑ appear as typical files in the file system
  - ❑ Created with the system call `mkfifo()` and manipulated with `open()`, `read()`, `write()` and `close()` system calls
  - ❑ Only half-duplex transmission is permitted; 2 FIFOs are used for data travel in both directions
  - ❑ Allows only byte-oriented data transmission across a FIFO
  - ❑ Communicating processes must reside on the same machine

- ❑ On Windows systems:
  - ❑ Created with the function CreateNamedPipe() and manipulated with ConnectNamedPipe(), ReadFile(), WriteFile() and DisconnectNamedPipe() functions
  - ❑ Full duplex transmission is permitted
  - ❑ Communicating processes may reside on the same or different machines
  - ❑ Allows byte- or message-oriented data transmission



# THANK YOU

---

**Chandravva Hebbi**

Department of Computer Science Engineering

**[chandravvahebbi@pes.edu](mailto:chandravvahebbi@pes.edu)**