



# OPERATING SYSTEMS

## Mutual Exclusion and Synchronization

---

**Chandravva Hebbi**

Department of Computer Science

# OPERATING SYSTEMS

---

## Software Approaches

**Chandravva Hebbi**

Department of Computer Science

# OPERATING SYSTEMS

## Slides Credits for all PPTs of this course

---



- The slides/diagrams in this course are an **adaptation**, **combination**, and **enhancement** of material from the following resources and persons:
  1. Slides of Operating System Concepts, Abraham Silberschatz, Peter Baer Galvin, Greg Gagne - 9<sup>th</sup> edition 2013 and some slides from 10<sup>th</sup> edition 2018
  2. Some conceptual text and diagram from Operating Systems - Internals and Design Principles, William Stallings, 9<sup>th</sup> edition 2018
  3. Some presentation transcripts from A. Frank – P. Weisberg
  4. Some conceptual text from Operating Systems: Three Easy Pieces, Remzi Arpaci-Dusseau, Andrea Arpaci Dusseau

- ❑ Processes can execute concurrently
  - ❑ May be interrupted at any time, partially completing execution
- ❑ Concurrent access to shared data may result in data inconsistency
- ❑ Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- ❑ Illustration of the problem:

Suppose that we wanted to provide a solution to the consumer-producer problem that fills *all* the buffers. We can do so by having an integer **counter** that keeps track of the number of full buffers.

Initially, **counter** is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.

```
while (true) {  
    /* produce an item in next_produced */  
  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```

- **counter++** could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

- **counter--** could be implemented as

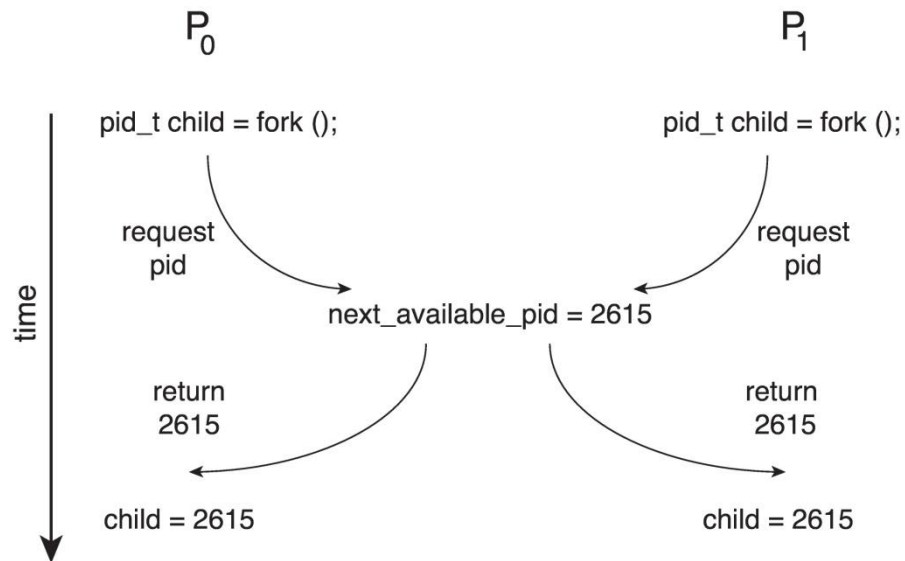
```
register2 = counter
register2 = register2 - 1
counter = register2
```

- Consider this execution interleaving with “count = 5” initially:

S0: producer execute	<b>register1 = counter</b>	{register1 = 5}
S1: producer execute	<b>register1 = register1 + 1</b>	{register1 = 6}
S2: consumer execute	<b>register2 = counter</b>	{register2 = 5}
S3: consumer execute	<b>register2 = register2 - 1</b>	{register2 = 4}
S4: producer execute	<b>counter = register1</b>	{counter = 6 }
S5: consumer execute	<b>counter = register2</b>	{counter = 4}

## Race Condition (Another Example)

- Processes  $P_0$  and  $P_1$  are creating child processes using the `fork()` system call
- Race condition on kernel variable `next_available_pid` which represents the next available process identifier (pid)



- Unless there is a mechanism to prevent  $P_0$  and  $P_1$  from accessing the variable `next_available_pid` the same pid could be assigned to two different processes!




- Consider system of  $n$  processes  $\{p_0, p_1, \dots, p_{n-1}\}$
- Each process has **critical section** segment of code
  - Process may be changing common variables, updating table, writing file, etc
  - When one process in critical section, no other may be in its critical section
- **Critical section problem** is to design protocol to solve this
- Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```

### □ General structure of process $P_i$

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```

1. **Mutual Exclusion** - If process  $P_i$  is executing in its critical section, then no other processes can be executing in their critical sections



2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
  - Assume that each process executes at a nonzero speed
  - No assumption concerning **relative speed** of the  $n$  processes

Two approaches depending on if kernel is preemptive or non-preemptive

- ❑ **Preemptive** – allows preemption of process when running in kernel mode.
- ❑ Preemptive kernels are difficult to design on SMP architectures
- ❑ **Non-preemptive** – runs until exits kernel mode, blocks, or voluntarily yields CPU
  - ▶ Essentially free of race conditions in kernel mode
  - ▶ It is free from race conditions on kernel data structures
- ❑ Preemptive kernel may be more responsive. Suitable for real-time systems.

- ❑ Software-based solution to the Critical Section problem.
- ❑ Good algorithmic description of solving the problem
- ❑ Assume that the **load** and **store** machine-language instructions are atomic; that is, cannot be interrupted
- ❑ **Peterson's Solution** restricted to two process.
- ❑  $P_i$  and  $P_j$  are two process,  $j=1-i$
- ❑ Peterson solution requires two processes share two data items:
  - ❑ `int turn;`
  - ❑ `Boolean flag[2]`
- ❑ The variable **turn** indicates whose turn it is to enter the critical section
- ❑ The **flag** array is used to indicate if a process is ready to enter the critical section. **flag[i] = true** implies that process  $P_i$  is ready!

```
while (true) {  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j);  
    /* critical section */  
    flag[i] = false;  
    /* remainder section */  
}
```

The structure of process  $P_i$  in Peterson's solution

- To prove solution is correct, we need show
- Mutual exclusion is preserved

$P_i$  enters critical section only if:

**turn = i**

and **turn** cannot be both 0 and 1 at the same time

- What about the Progress requirement?
- What about the Bounded-waiting requirement?



□ Provable that the three CS requirement are met:

1. Mutual exclusion is preserved

$P_i$  enters CS only if:

either **flag[j] = false** or **turn = i**

2. Progress requirement is satisfied

3. Bounded-waiting requirement is met

To prove properties 2 and 3

- we note that a process  $P_i$  can be prevented from entering the critical section only if it is stuck in the while loop with the condition
  - $\text{flag}[j] == \text{true}$  and  $\text{turn} == j$ ; this loop is the only one possible.
- If  $P_j$  is not ready to enter the critical section, then  $\text{flag}[j] == \text{false}$ , and  $P_i$  can enter its critical section.
- If  $P_j$  has set  $\text{flag}[j]$  to true and is also executing in its while statement, then either  $\text{turn} == i$  or  $\text{turn} == j$ . If  $\text{turn} == i$ , then  $P_i$  will enter the critical section.
- If  $\text{turn} == j$ , then  $P_j$  will enter the critical section.
- once  $P_j$  exits its critical section, it will reset  $\text{flag}[j]$  to false, allowing  $P_i$  to enter its critical section.
- If  $P_j$  resets  $\text{flag}[j]$  to true, it must also set  $\text{turn}$  to  $i$ .
- Thus, since  $P_i$  does not change the value of the variable  $\text{turn}$  while executing the while statement.
- $P_i$  will enter the critical section (progress) after at most one entry by  $P_j$  (bounded waiting).



# THANK YOU

---

**Chandravva Hebbi**

Department of Computer Science Engineering

**[chandravvahebbs@pes.edu](mailto:chandravvahebbs@pes.edu)**