

# OPERATING SYSTEMS

---

## Memory Management

**Chandravva Hebbi**

Department of Computer Science

# OPERATING SYSTEMS

---

**TLBs context switches**

**Chandravva Hebbi**

Department of Computer Science

- The slides/diagrams in this course are an **adaptation**, **combination**, and **enhancement** of material from the following resources and persons:
  1. Slides of Operating System Concepts, Abraham Silberschatz, Peter Baer Galvin, Greg Gagne - 9<sup>th</sup> edition 2013 and some slides from 10<sup>th</sup> edition 2018
  2. Some conceptual text and diagram from Operating Systems - Internals and Design Principles, William Stallings, 9<sup>th</sup> edition 2018
  3. Some presentation transcripts from A. Frank – P. Weisberg
  4. Some conceptual text from Operating Systems: Three Easy Pieces, Remzi Arpaci-Dusseau, Andrea Arpaci Dusseau

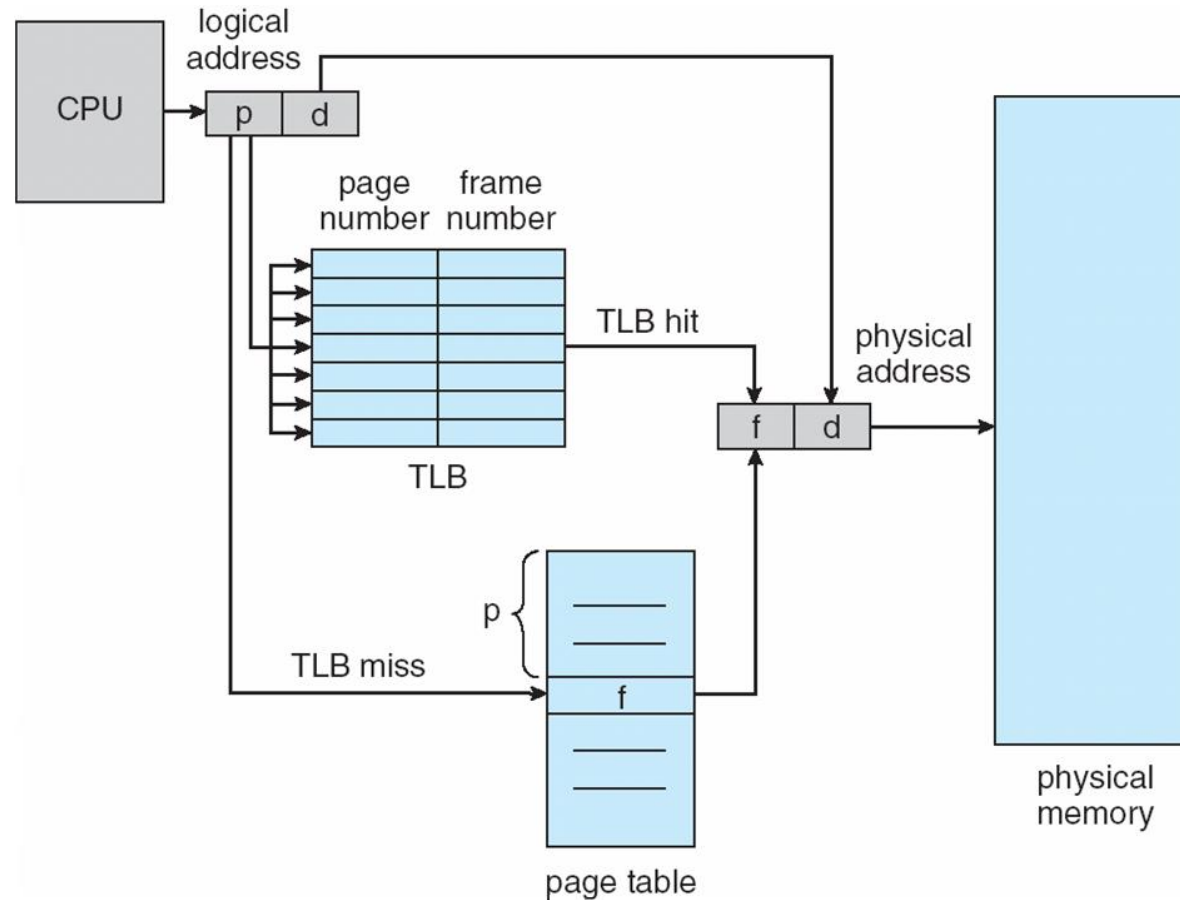
- ❑ The hardware implementation of page table can be done by using dedicated registers.
- ❑ But the usage of register for the page table is satisfactory only if page table is small.
- ❑ If page table contain large number of entries then we can use TLB(translation Look-aside buffer), a special, small, fast look up hardware cache.

- ❑ Page table is kept in main memory
- ❑ **Page-table base register (PTBR)** points to the page table
- ❑ **Page-table length register (PTLR)** indicates size of the page table
- ❑ In this scheme every data/instruction access requires two memory accesses
  - ❑ One for the page table and one for the data / instruction
- ❑ The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**

- ❑ Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
  - ❑ Otherwise need to flush at every context switch
- ❑ TLBs typically small (64 to 1,024 entries)
- ❑ On a TLB miss, value is loaded into the TLB for faster access next time
  - ❑ Replacement policies must be considered
  - ❑ Some entries can be **wired down** for permanent fast access

# OPERATING SYSTEMS

## Paging Hardware With TLB



- Hit ratio – percentage of times that a page number is found in the TLB
- An 80% hit ratio means that we find the desired page number in the TLB 80% of the time.
- Suppose that 10 nanoseconds to access memory.
  - If we find the desired page in TLB then a mapped-memory access take 10 ns
  - Otherwise we need two memory access so it is 20 ns i.e. if we fail to find the page number in the TLB then we must first access memory for the page table and frame number (10 ns) and then access the desired byte in memory (10 nanoseconds), for a total of 20 ns (assuming that a page table lookup takes only one memory access).

- **Effective Access Time (EAT)**

$$\text{EAT} = 0.80 \times 10 + 0.20 \times 20 = 12 \text{ nanoseconds}$$

implying 20% slowdown in access time

- Consider a more realistic hit ratio of 99%,

$$\text{EAT} = 0.99 \times 10 + 0.01 \times 20 = 10.1\text{ns}$$

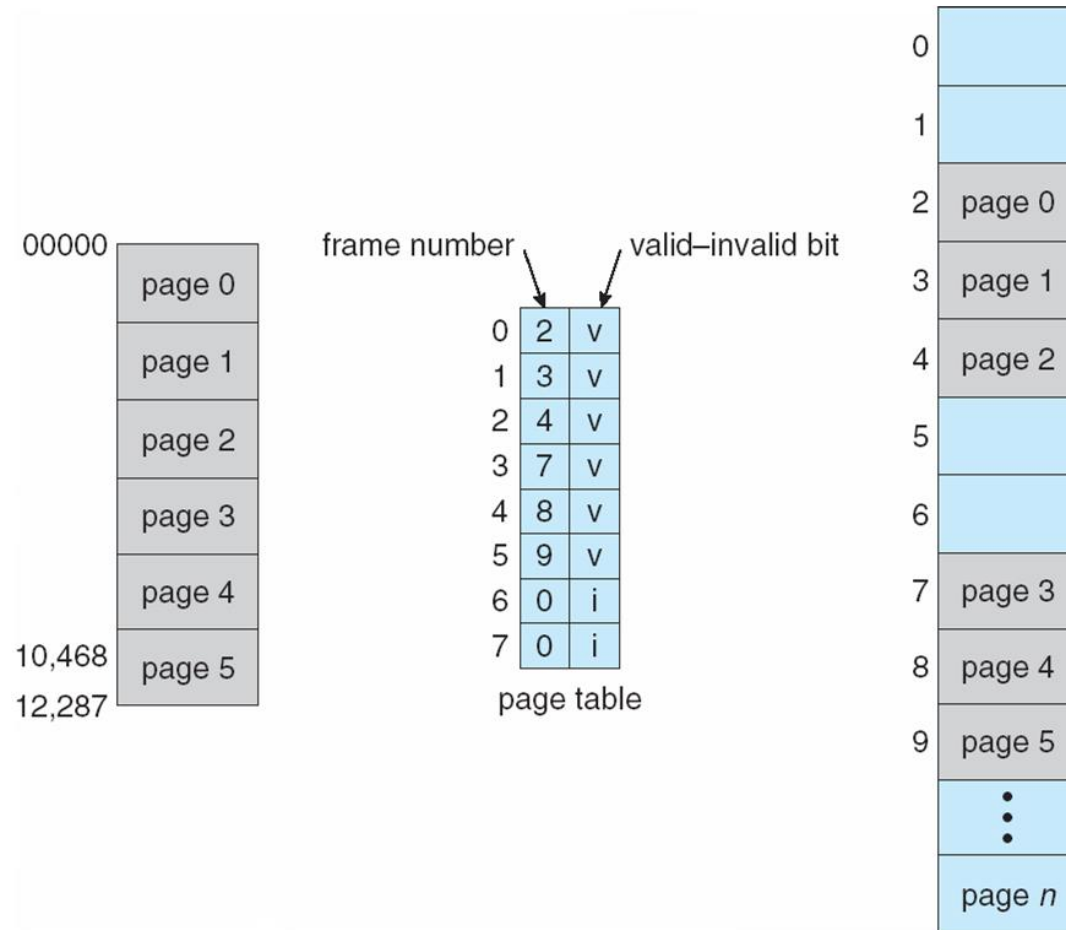
implying only 1% slowdown in access time.



- ❑ Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
  - ❑ Can also add more bits to indicate page execute-only, and so on
- ❑ **Valid-invalid** bit attached to each entry in the page table:
  - ❑ “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
  - ❑ “invalid” indicates that the page is not in the process’ logical address space
  - ❑ Or use **page-table length register (PTLR)**
- ❑ Any violations result in a trap to the kernel

# OPERATING SYSTEMS

## Valid (v) or Invalid (i) Bit In A Page Table



### □ Shared code

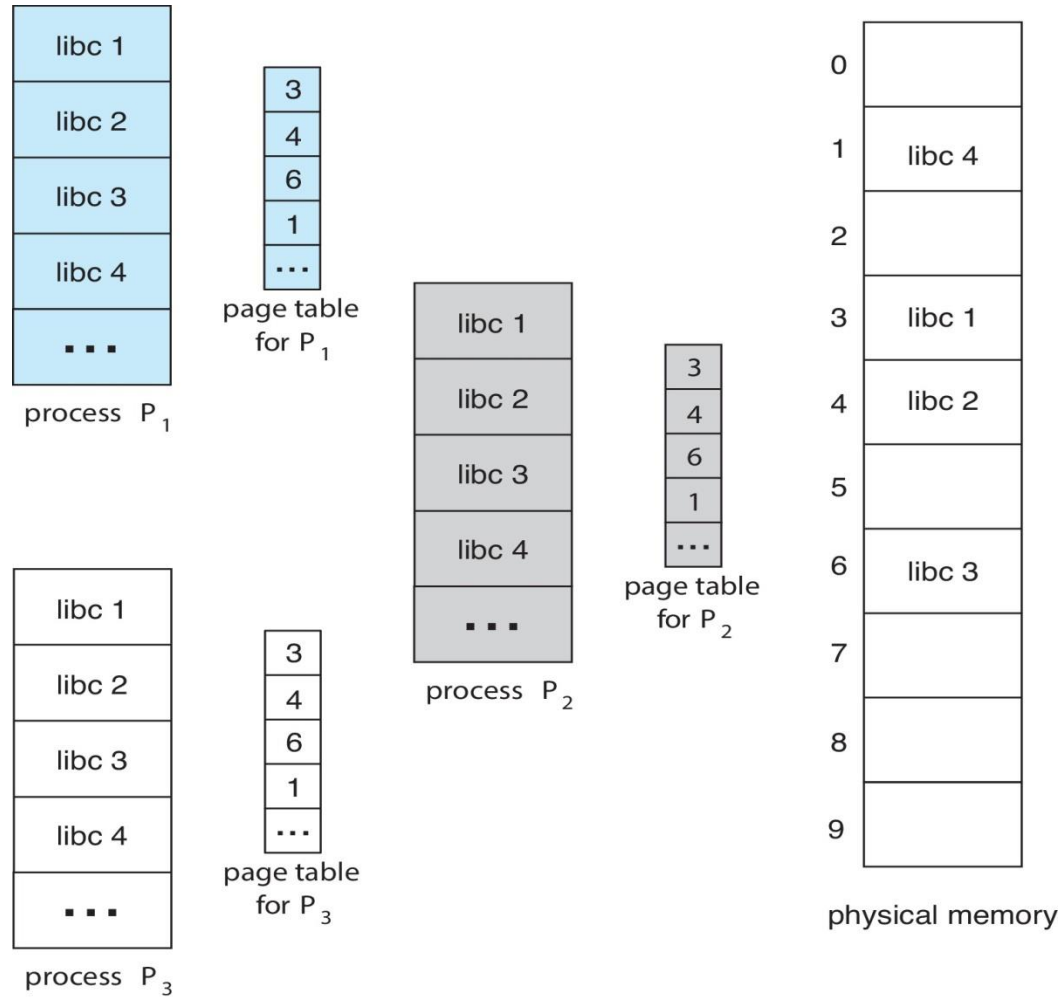
- One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems)
- Similar to multiple threads sharing the same process space
- Also useful for interprocess communication if sharing of read-write pages is allowed

### □ Private code and data

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space

# OPERATING SYSTEMS

## Shared Pages Example





# THANK YOU

---

**Chandravva Hebbi**

Department of Computer Science Engineering

**[chandravvahebbi@pes.edu](mailto:chandravvahebbi@pes.edu)**