

# OPERATING SYSTEMS

---

## Memory Management

**Chandravva Hebi**

Department of Computer Science

# OPERATING SYSTEMS

---

## Virtual Memory

**Chandravva Hebi**

Department of Computer Science

- The slides/diagrams in this course are an **adaptation**, **combination**, and **enhancement** of material from the following resources and persons:
  1. Slides of Operating System Concepts, Abraham Silberschatz, Peter Baer Galvin, Greg Gagne - 9<sup>th</sup> edition 2013 and some slides from 10<sup>th</sup> edition 2018
  2. Some conceptual text and diagram from Operating Systems - Internals and Design Principles, William Stallings, 9<sup>th</sup> edition 2018
  3. Some presentation transcripts from A. Frank – P. Weisberg
  4. Some conceptual text from Operating Systems: Three Easy Pieces, Remzi Arpaci-Dusseau, Andrea Arpaci Dusseau

# OPERATING SYSTEMS

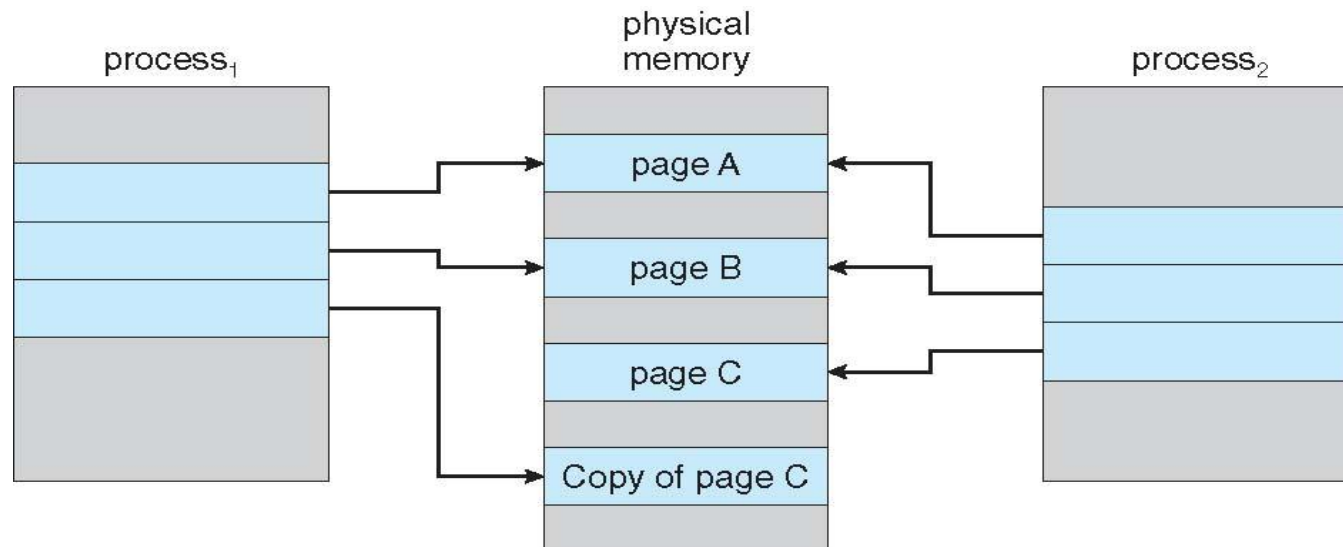

## Copy-on-Write



- ❑ **Copy-on-Write** (COW) allows both parent and child processes to initially *share* the same pages in memory
  - ❑ If either process modifies a shared page, only then is the page copied
- ❑ COW allows more efficient process creation as only modified pages are copied
- ❑ In general, free pages are allocated from a **pool** of **zero-fill-on-demand** pages
  - ❑ Pool should always have free frames for fast demand page execution
    - ▶ Don't want to have to free a frame as well as other processing on page fault
  - ❑ Why zero-out a page before allocating it?
- ❑ `vfork()` variation on `fork()` system call has parent suspend and child using copy-on-write address space of parent
  - ❑ Designed to have child call `exec()`
  - ❑ Very efficient

# OPERATING SYSTEMS

## Before and After Process 1 Modifies Page C



## What Happens if There is no Free Frame?


---

- ❑ Used up by process pages
- ❑ Also in demand from the kernel, I/O buffers, etc
- ❑ How much to allocate to each?
- ❑ Page replacement – find some page in memory, but not really in use, page it out
  - ❑ Algorithm – terminate? swap out? replace the page?
  - ❑ Performance – want an algorithm which will result in minimum number of page faults
- ❑ Same page may be brought into memory several times

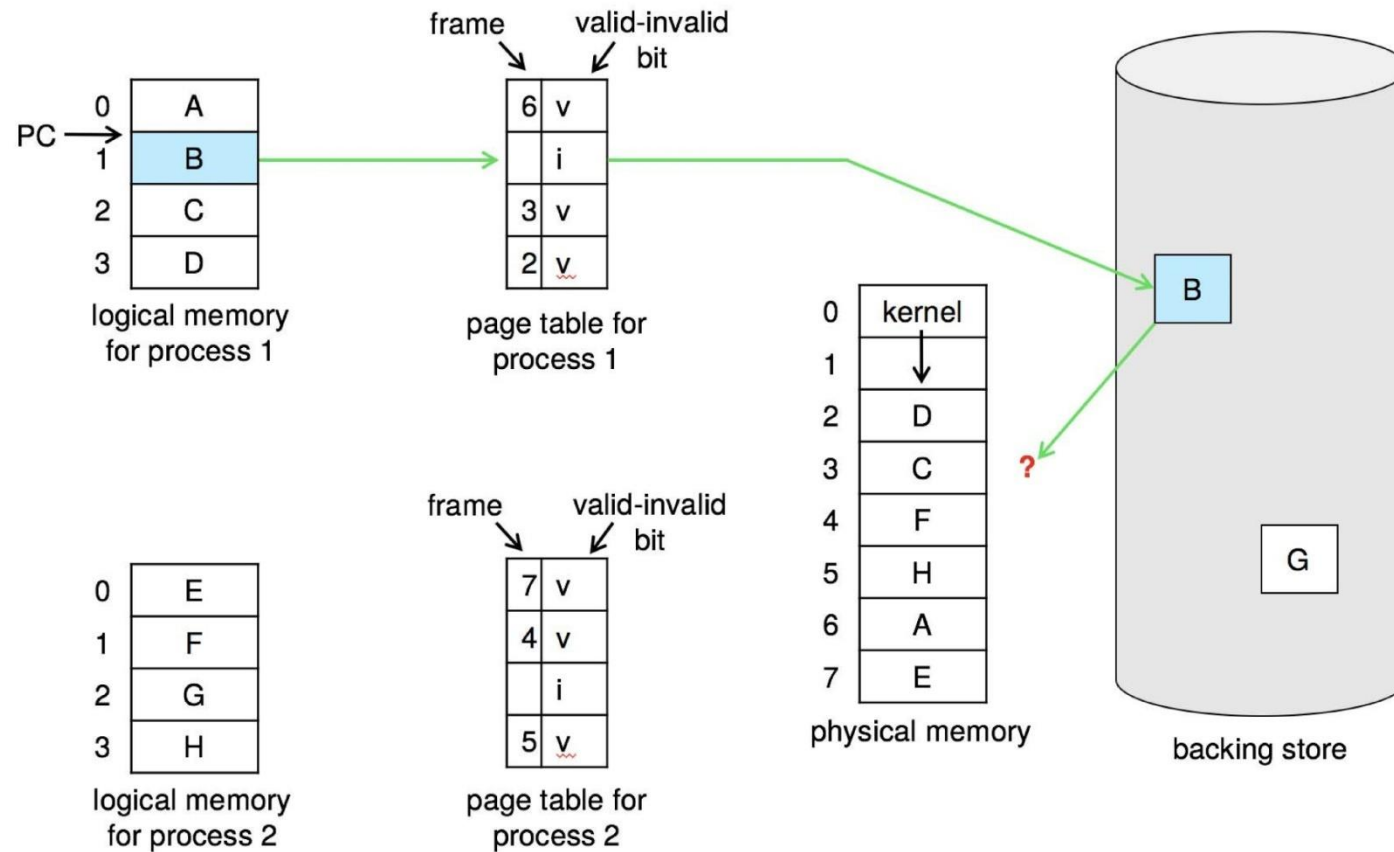
- ❑ Prevent **over-allocation** of memory by modifying page-fault service routine to include page replacement (vs simply increasing degree of multiprogramming)
- ❑ Use **modify (dirty) bit** to reduce overhead of page transfers – only modified pages are written to disk
- ❑ Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory

# OPERATING SYSTEMS

## Need For Page Replacement – Example 1








1. Find the location of the desired page on disk
2. Find a free frame:
  - If there is a free frame, use it
  - If there is no free frame, use a page replacement algorithm to select a **victim frame**
    - Write victim frame to disk if dirty
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Continue the process by restarting the instruction that caused the trap

Note now potentially 2 page transfers for page fault – increasing EAT

# OPERATING SYSTEMS

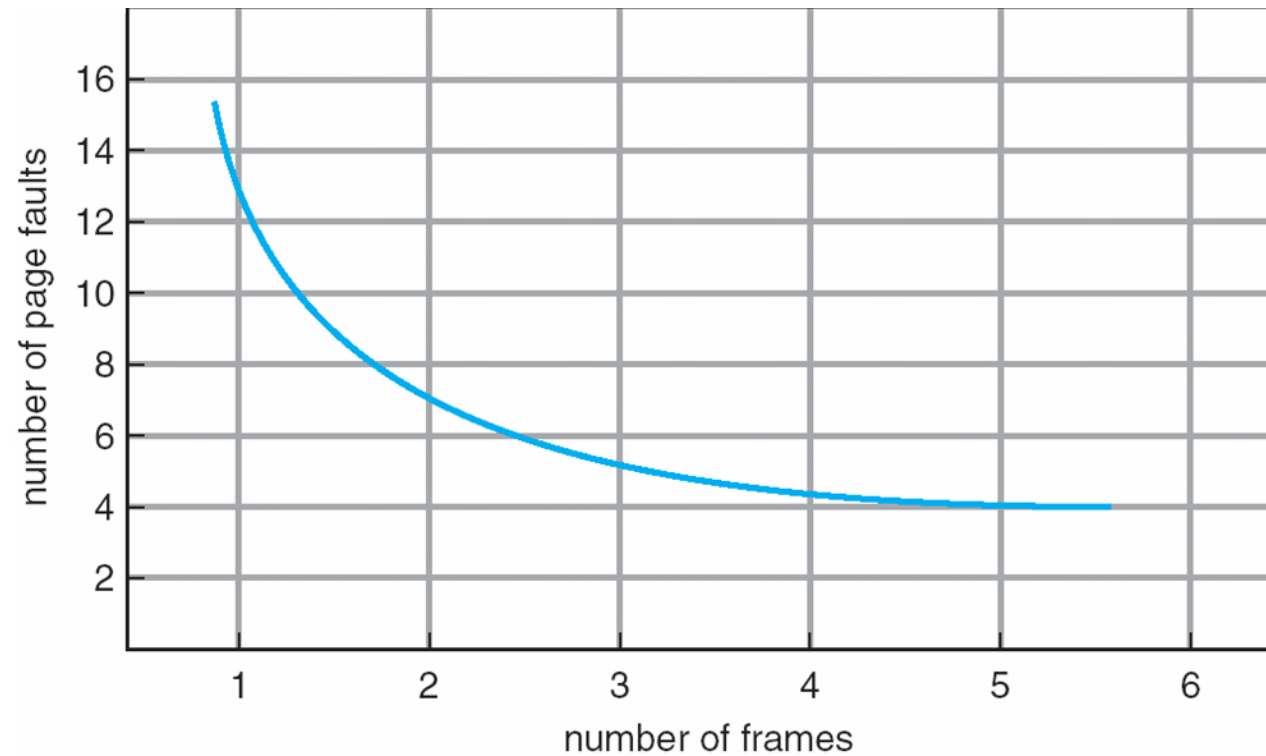
## Basic Page Replacement (Cont.)



- ❑ **Frame-allocation algorithm** determines
  - ❑ How many frames to give each process
  - ❑ Which frames to replace
- ❑ **Page-replacement algorithm**
  - ❑ Want lowest page-fault rate on both first access and re-access
- ❑ Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
  - ❑ String is just page numbers, not full addresses
  - ❑ Repeated access to the same page does not cause a page fault
  - ❑ Results depend on number of frames available
- ❑ In all our examples, the **reference string** of referenced page numbers is

**7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1**

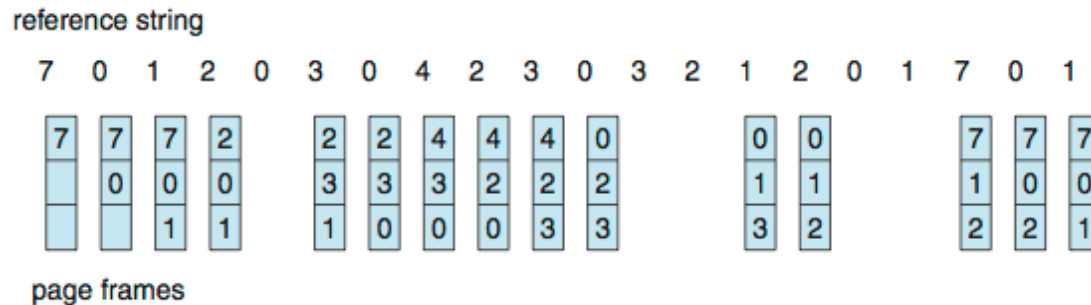
## Graph of Page Faults Versus The Number of Frames




# OPERATING SYSTEMS

## First-In-First-Out (FIFO) Algorithm

- Reference string: **7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1**
- 3 frames (3 pages can be in memory at a time per process)



- Can vary by reference string: consider 1,2,3,4,1,2,5,1,2,3,4,5
  - Adding more frames can cause more page faults!
  - **Belady's Anomaly**
- How to track ages of pages?
  - Just use a FIFO queue



# OPERATING SYSTEMS

## Optimal Algorithm

- Replace page that will not be used for longest period of time
  - 9 is optimal for the example
- How do you know this?
  - Can't read the future
- Used for measuring how well your algorithm performs

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		2		2		2						7		
	0	0	0		0		4		0		0						0		
		1	1		3		3		3		1						1		

page frames



# OPERATING SYSTEMS

## Least Recently Used (LRU) Algorithm

- Use past knowledge rather than future
- Replace page that has not been used in the most amount of time
- Associate time of last use with each page

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames

- 12 faults – better than FIFO but worse than OPT
- Generally good algorithm and frequently used
- But how to implement?

# OPERATING SYSTEMS

## LRU Algorithm (Cont.)

---




- ❑ Counter implementation
  - ❑ Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
  - ❑ When a page needs to be changed, look at the counters to find smallest value
    - ▶ Search through table needed
- ❑ Stack implementation
  - ❑ Keep a stack of page numbers in a double link form:
  - ❑ Page referenced:
    - ▶ move it to the top
    - ▶ requires 6 pointers to be changed
  - ❑ But each update more expensive
  - ❑ No search for replacement

- LRU and OPT are cases of **stack algorithms** that don't have Belady's Anomaly
- Use of a Stack to Record Most Recent Page References

reference string


4 7 0 7 1 0 1 2 1 2 7 1 2



stack  
before  
a



stack  
after  
b





**THANK YOU**

---

**Chandravva Hebi**

Department of Computer Science Engineering

**[chandravvahebbi@pes.edu](mailto:chandravvahebbi@pes.edu)**