

OPERATING SYSTEMS

Memory Management

Chandravva Hebi

Department of Computer Science

OPERATING SYSTEMS

Virtual Memory

Chandravva Hebi

Department of Computer Science

- The slides/diagrams in this course are an **adaptation**, **combination**, and **enhancement** of material from the following resources and persons:
 1. Slides of Operating System Concepts, Abraham Silberschatz, Peter Baer Galvin, Greg Gagne - 9th edition 2013 and some slides from 10th edition 2018
 2. Some conceptual text and diagram from Operating Systems - Internals and Design Principles, William Stallings, 9th edition 2018
 3. Some presentation transcripts from A. Frank – P. Weisberg
 4. Some conceptual text from Operating Systems: Three Easy Pieces, Remzi Arpaci-Dusseau, Andrea Arpaci Dusseau

OPERATING SYSTEMS

Least Recently Used (LRU) Algorithm

- Use past knowledge rather than future
- Replace page that has not been used in the most amount of time
- Associate time of last use with each page

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames

- 12 faults – better than FIFO but worse than OPT
- Generally good algorithm and frequently used
- But how to implement?

- ❑ Counter implementation
 - ❑ Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
 - ❑ When a page needs to be changed, look at the counters to find smallest value
 - ▶ Search through table needed
- ❑ Stack implementation
 - ❑ Keep a stack of page numbers in a double link form:
 - ❑ Page referenced:
 - ▶ move it to the top
 - ▶ requires 6 pointers to be changed
 - ❑ But each update more expensive
 - ❑ No search for replacement
- ❑ LRU and OPT are cases of **stack algorithms** that don't have Belady's Anomaly

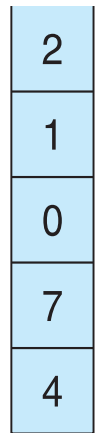
OPERATING SYSTEMS

Use Of A Stack to Record Most Recent Page References

- LRU and OPT are cases of **stack algorithms** that don't have Belady's Anomaly
- Use Of A Stack to Record Most Recent Page References

reference string

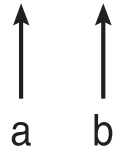
4 7 0 7 1 0 1 2 1 2 7 1 2



stack
before
a



stack
after
b



- ❑ Each process needs ***minimum*** number of frames
- ❑ Example: IBM 370 – 6 pages to handle SS MOVE instruction:
 - ❑ instruction is 6 bytes, might span 2 pages
 - ❑ 2 pages to handle *from*
 - ❑ 2 pages to handle *to*
- ❑ ***Maximum*** of course is total frames in the system
- ❑ Two major allocation schemes
 - ❑ fixed allocation
 - ❑ priority allocation
- ❑ Many variations

- Equal allocation – For example, if there are 100 frames (after allocating frames for the OS) and 5 processes, give each process 20 frames
 - Keep some as free frame buffer pool
- Proportional allocation – Allocate according to the size of process
 - Dynamic as degree of multiprogramming, process sizes change

- s_i = size of process p_i
- $S = \sum s_i$
- m = total number of frames
- a_i = allocation for $p_i = \frac{s_i}{S} \times m$

$$m = 62$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 62 \approx 4$$

$$a_2 = \frac{127}{137} \times 62 \approx 57$$

- Use a proportional allocation scheme using priorities rather than size

- If process P_i generates a page fault,
 - select for replacement one of its frames
 - select for replacement a frame from a process with lower priority number

- ❑ **Global replacement** – process selects a replacement frame from the set of all frames; one process can take a frame from another
 - ❑ But then process execution time can vary greatly
 - ❑ But greater throughput so more common

- ❑ **Local replacement** – each process selects from only its own set of allocated frames
 - ❑ More consistent per-process performance
 - ❑ But possibly underutilized memory

- ❑ So far all memory accessed equally
- ❑ Many systems are **NUMA** – speed of access to memory varies
 - ❑ Consider system boards containing CPUs and memory, interconnected over a system bus
- ❑ Optimal performance comes from allocating memory “close to” the CPU on which the thread is scheduled
 - ❑ And modifying the scheduler to schedule the thread on the same system board when possible
 - ❑ Solved by Solaris by creating **lgroups**
 - ▶ Structure to track CPU / Memory low **latency groups**
 - ▶ Used my schedule and pager
 - ▶ When possible schedule all threads of a process and allocate all memory for that process within the lgroup (otherwise it picks nearby lgroups)



THANK YOU

Chandravva Hebi

Department of Computer Science Engineering

chandravvahebbi@pes.edu