

# OPERATING SYSTEMS

---

## I/O Management, System Protection and Security

**Kakoli Bora**

Department of Computer Science

# OPERATING SYSTEMS

---

**I/O Hardware, polling and interrupts**

**Kakoli Bora**

Department of Computer Science

# OPERATING SYSTEMS

## Slides Credits for all PPTs of this course

---



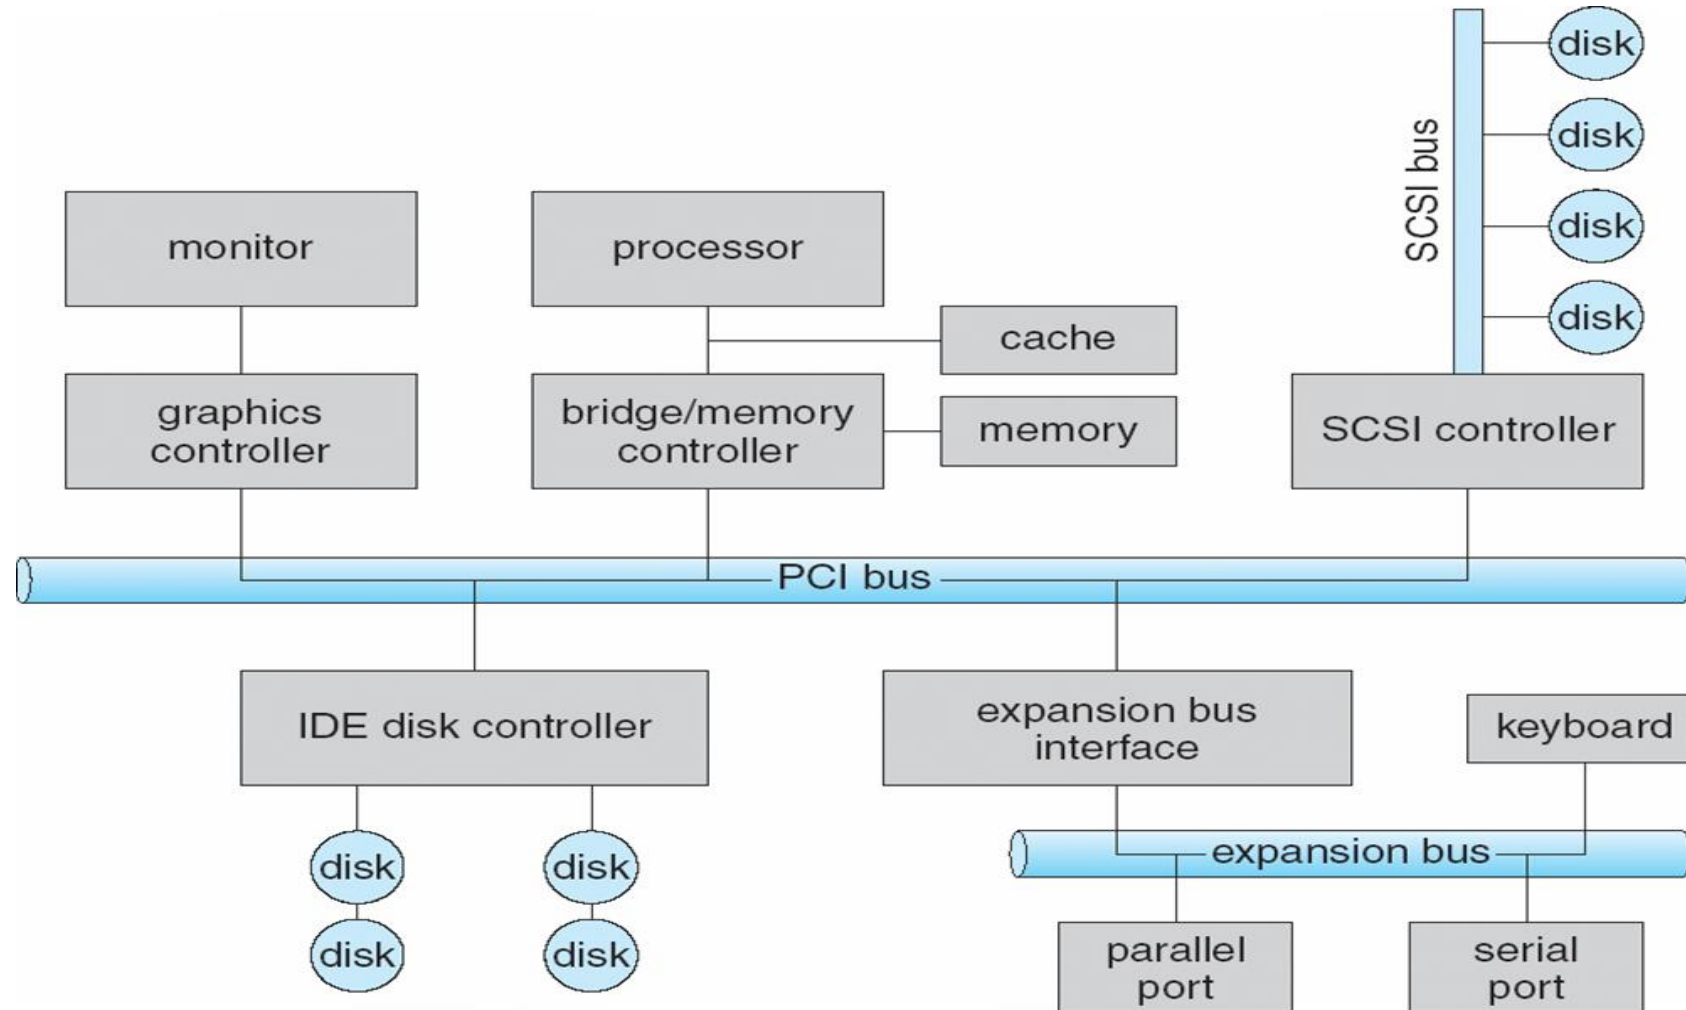
- The slides/diagrams in this course are an **adaptation, combination,** and **enhancement** of material from the following resources and persons:
1. Slides of Operating System Concepts, Abraham Silberschatz, Peter Baer Galvin, Greg Gagne - 9<sup>th</sup> edition 2013 and some slides from 10<sup>th</sup> edition 2018
  2. Some conceptual text and diagram from Operating Systems - Internals and Design Principles, William Stallings, 9<sup>th</sup> edition 2018
  3. Some presentation transcripts from A. Frank – P. Weisberg
  4. Some conceptual text from Operating Systems: Three Easy Pieces, Remzi Arpaci-Dusseau, Andrea Arpaci Dusseau

- ❑ I/O management is a major component of operating system design and operation
  - ❑ Important aspect of computer operation
  - ❑ I/O devices vary greatly
  - ❑ Various methods to control them
  - ❑ Performance management
  - ❑ New types of devices frequent
- ❑ Ports, buses, device controllers connect to various devices
- ❑ **Device drivers** encapsulate device details
  - ❑ Present uniform device-access interface to I/O subsystem

- ❑ Incredible variety of I/O devices
  - ❑ Storage (Disks, tapes)
  - ❑ Transmission (network connections, Bluetooth)
  - ❑ Human-interface (screen, keyboard, mouse, audio in and out)
- ❑ Common concepts – signals from I/O devices interface with computer
  - ❑ **Port** – connection point for device
  - ❑ **Bus - daisy chain** or shared direct access
    - ▶ **PCI** bus common in PCs and servers, PCI Express (**PCIe**)
    - ▶ **expansion bus** connects relatively slow devices
  - ❑ **Controller (host adapter)** – electronics that operate port, bus, device
    - ▶ Sometimes integrated
    - ▶ Sometimes separate circuit board (host adapter)
    - ▶ Contains processor, microcode, private memory, bus controller, etc
      - Some talk to per-device controller with bus controller, microcode, memory, etc

# OPERATING SYSTEMS

## A Typical PC Bus Structure



- I/O instructions control devices
- Devices usually have registers where device driver places commands, addresses, and data to write, or read data from registers after command execution
  - Data-in register, data-out register, status register, control register
  - Typically 1-4 bytes, or FIFO buffer

- Devices have addresses, used by
  - **Direct I/O instructions**
    - ▶ Special I/O instructions specify the transfer of a byte or word to an I/O port address.
    - ▶ The I/O instruction triggers bus lines to select the proper device and to move bits into or out of a device register.
  - **Memory-mapped I/O**
    - ▶ Device data and command registers mapped to processor address space
    - ▶ Especially for large address spaces (graphics)
- Some systems use both techniques – I/O instructions to control some devices and memory-mapped I/O to control others



# OPERATING SYSTEMS

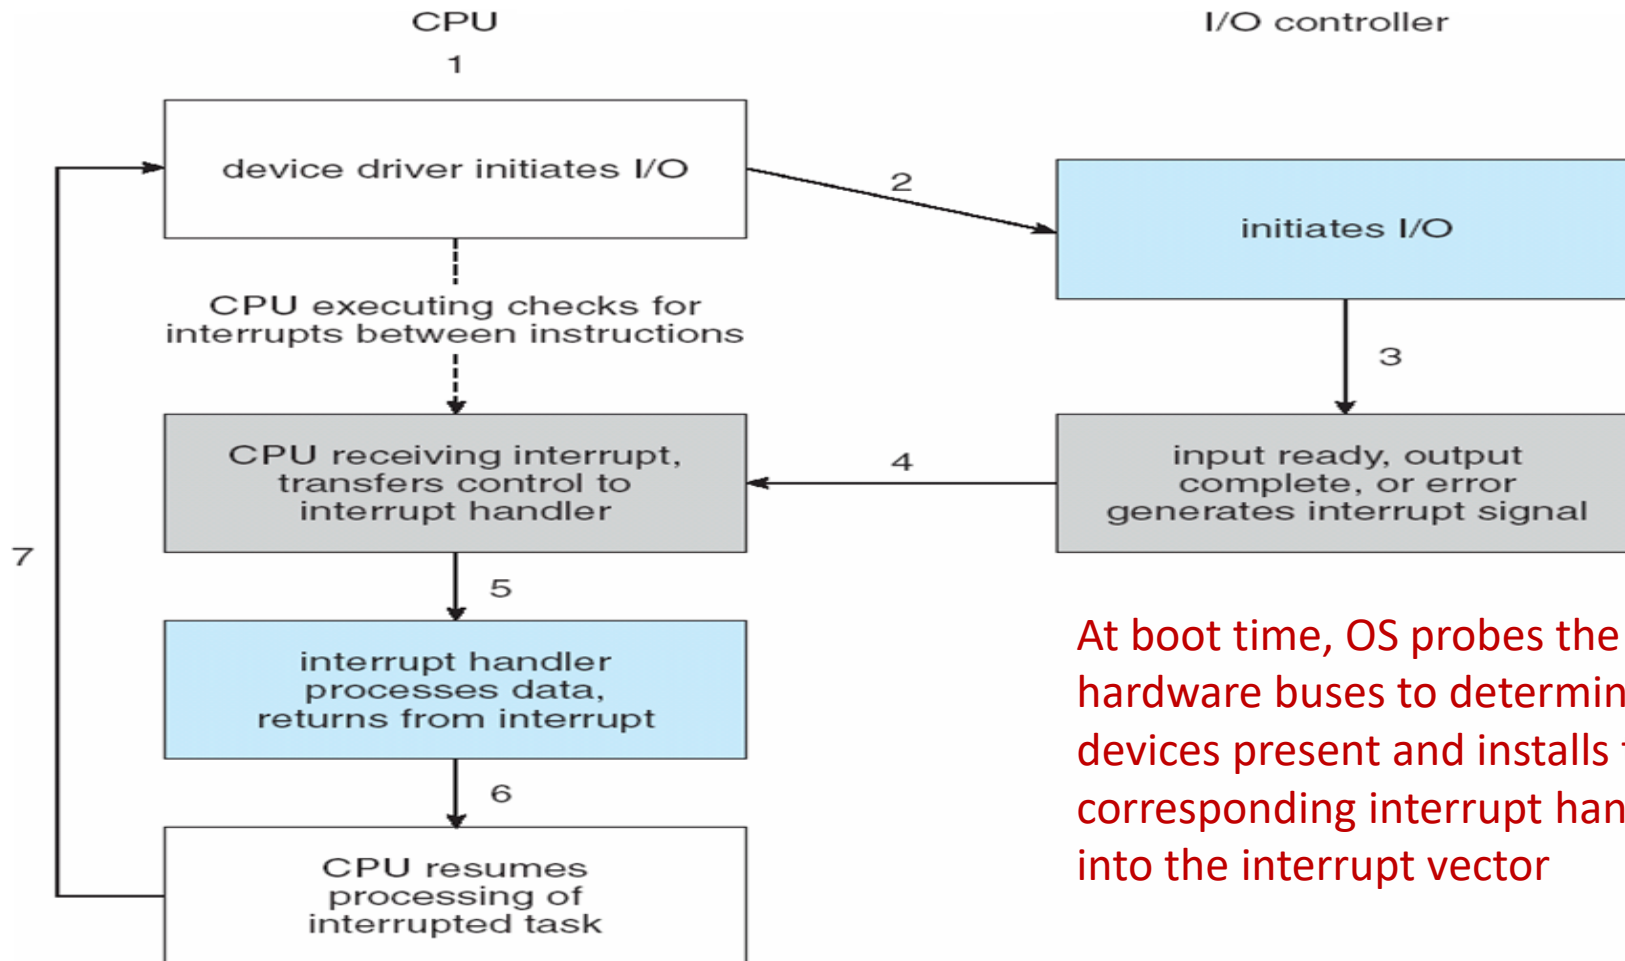
## Device I/O Port Locations on PCs (partial)

---

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)

- For each byte of I/O
  1. Host repeatedly reads busy bit from status register until that bit becomes clear (**here host is busy-waiting or polling**)
  2. Host sets read or write bit and if write, copies data into data-out register
  3. Host sets command-ready bit
  4. Controller notices command-ready bit is set, sets busy bit, executes transfer
  5. Controller clears busy bit, error bit, command-ready bit when transfer done
- Step 1 above is **busy-wait** cycle to wait for I/O from device
  - Reasonable if device is fast
  - But inefficient if device slow
  - CPU switches to other tasks?
    - ▶ But if miss a cycle data overwritten / lost

- ❑ Polling can happen in 3 instruction cycles
  - ❑ Read a device register, logical--and to extract status bit, branch if not zero
  - ❑ How to be more efficient if non-zero infrequently?
- ❑ CPU **Interrupt-request line** triggered by I/O device
  - ❑ Checked by processor after each instruction
- ❑ **Interrupt handler** receives interrupts
  - ❑ **Maskable** to ignore or delay some interrupts
- ❑ **Interrupt vector** to dispatch interrupt to correct handler
  - ❑ Context switch at start and end
  - ❑ Based on priority
  - ❑ Some **nonmaskable**
  - ❑ Interrupt chaining (list of interrupt handlers) if more than one device at same interrupt number
- ❑ **Above features are provided by the CPU and by the Interrupt-controller hardware**
  - ❑ **Most CPUs have 2 interrupt request lines – maskable and nonmaskable**



At boot time, OS probes the hardware buses to determine the devices present and installs the corresponding interrupt handlers into the interrupt vector

vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts

Events from 0 to 31 are nonmaskable and are used to signal various error conditions

- ❑ Interrupt mechanism also used for **exceptions**
  - ❑ Terminate process, crash system due to hardware error
- ❑ Page fault executes when memory access error
- ❑ System call executes via **trap** to trigger kernel to execute request
  - Software interrupt or trap is given a lower priority compared to device interrupts
- ❑ Multi-CPU systems can process interrupts concurrently
  - ❑ If operating system designed to handle it
- ❑ Used for time-sensitive processing, frequent, must be fast
  - ❑ **Note: Interrupt handlers in Solaris are executed as kernel threads**



# THANK YOU

---

**Kakoli Bora**

Department of Computer Science Engineering

**[k\\_bora@pes.edu](mailto:k_bora@pes.edu)**